COMP5046
Natural Language Processing

*Lecture 10: Attention and Question Answering
(Reading Comprehension)*

Dr. Caren Han

*Semester 1, 2021
School of Computer Science,
University of Sydney*

Deep Learning

Linguistics

Language

NLP

LECTURE PLAN

Lecture 10: Attention and Question Answering (Reading Comprehension)

1. Question Answering
2. Knowledge-based Question Answering
3. IR-based Question Answering (Reading Comprehension)
4. Attention
5. Reading Comprehension with Attention
6. Visual Question Answering

Question Answering

Question Answering

Question answering (QA) *is a computer science discipline within the fields of information retrieval and natural language processing (NLP), which is concerned with building systems that* ***automatically answer questions posed by humans in a natural language***.

Different types of questions:

General questions, with Yes/No answers
- *e.g. Are you a student?*

Wh- Questions, start with: who, what, where, when, why, how, how many
- *e.g. When did you get to this lecture?*
- *e.g. What is the weather like in London?*

Question Answering

Question Answering

Question answering (QA) *is a computer science discipline within the fields of information retrieval and natural language processing (NLP), which is concerned with building systems that* ***automatically answer questions posed by humans in a natural language****.*

Different types of questions:

Choice Questions*, where you have some options inside the question*

Factoid questions*, where the complete answer can be found inside a text. The answer to such questions consist of one or several words that go one after another*

Question Answering

Question

Three Questions for building a QA System

- What does the argument look like?
- Where can I get the answers from?
- What does my training data look like?

THE UNIVERSITY OF
SYDNEY

Research Areas in Question Answering

Research Area	Details
Knowledge-based QA (Semantic Parsing)	• Answer is a logical form, possible executed against a Knowledge Base • Context is a Knowledge Base
Information Retrieval-based QA • Answer sentence selection • Reading Comprehension	• Answer is a document, paragraph, sentence • Context is a corpus of documents or a specific document
Visual QA	• Answer is simple and factual • Context is one/multiple image(s)
Library Reference	• Answer is another question • Context is the structured knowledge available in the library and the librarians view of it.

Knowledge-based Question Answering

Semantic Parsing

Answering a natural language question by mapping it to a query over a structured database (formal representation of its meaning).

Assignment Project Exam Help

Question

When was Justin Bieber born?

https://tutorcs.com

Logical Form

birth-year(Justin Bieber, x)

WeChat: cstutorcs

KB Query

Knowledge base

Answer

Justin Bieber was born in 1994

Knowledge-based Question Answering

Semantic Parsing

Answering a natural language question by mapping it to a query over a structured database (formal representation of its meaning).

Question

When was Justin Bieber born?

Logical Form

birth-year(Justin Bieber, x)

KB Query

Knowledge base

Answer

Justin Bieber was born in 1994
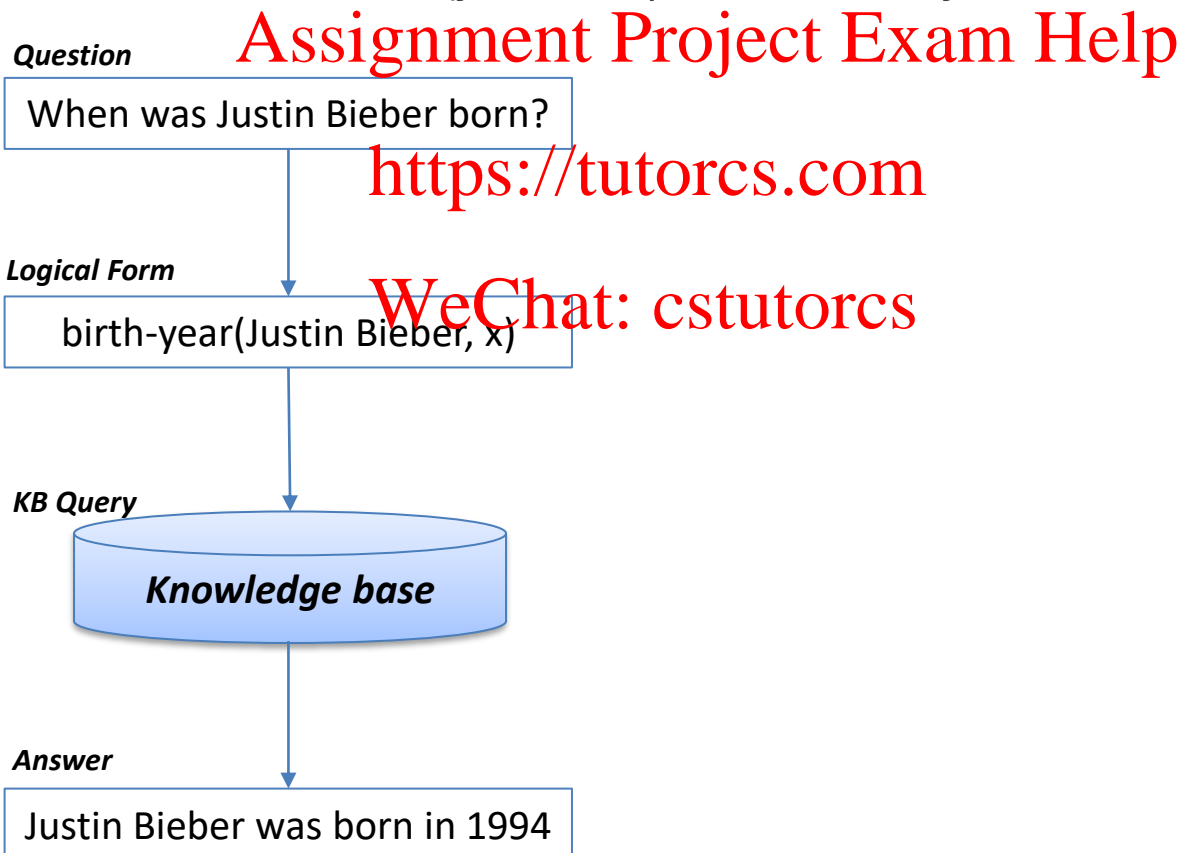
Mapping from a text string to any logical form

Question	Logical Form
When was Justin Bieber born?	birth-year(Justin Bieber, x)
What is the largest state?	argmax(λx.state(x),λx.size(x))

How to map? Map either to <u>some version of predicate calculus</u> or a query language like SQL or SPARQL
https://query.wikidata.org/

Use expensive supervised data?

Require experts for the manual annotation process....

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

Knowledge-based Question Answering

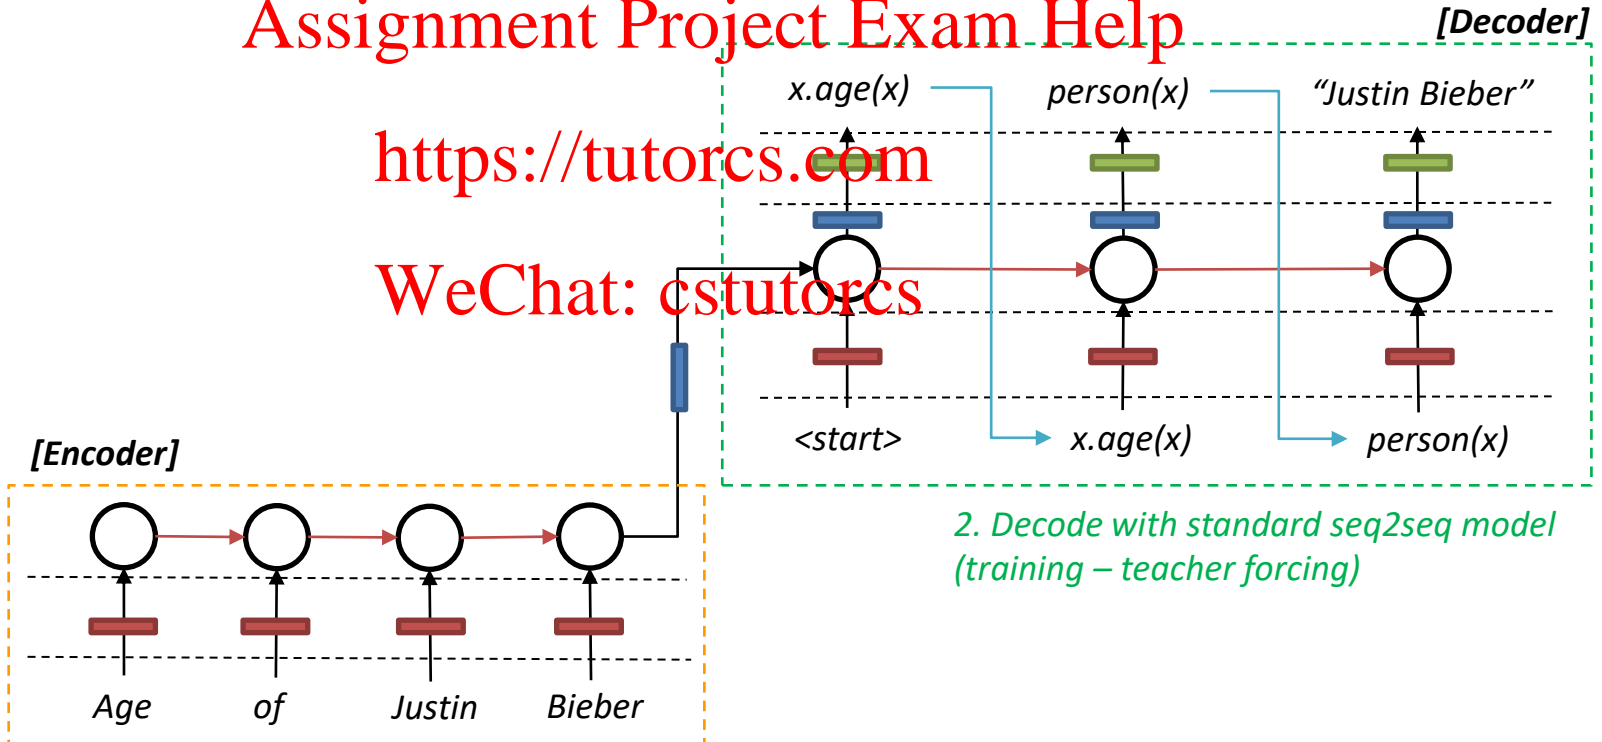THE UNIVERSITY OF SYDNEY

Seq2Seq model for semantic parser

How to transfer the text to the logical form?

A basic deep learning approach to semantic parsing

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

[Decoder]

x.age(x)   person(x)   "Justin Bieber"

<start>   x.age(x)   person(x)

[Encoder]

Age   of   Justin   Bieber

2. Decode with standard seq2seq model (training – teacher forcing)

1. Encode sentence with sequence models

Knowledge-based Question Answering

Semantic Parsing

Answering a natural language question by mapping it to a query over a structured database (formal representation of its meaning).

Question

When was Justin Bieber born?

Logical Form

birth-year(Justin Bieber, x)

KB Query

Knowledge base

Answer

Justin Bieber was born in 1994

Answer questions that ask about one of the missing arguments in a triple

Subject	Predicate (relation)	Object
Justin Bieber	birth-year	**1994**
Frédéric Chopin	birth-year	1810
…	…	…

- DBPedia
- Freebase

How to produce the answer?
- *Seq2seq*
- *Template based generation*

Knowledge-based Question Answering

Pros and Cons of Knowledge-based QA

- *Logical Form instead of (direct) answer makes system robust*
- *Answer independent of question and parsing mechanism*

- *Constrained to queriable questions in Database Schema*
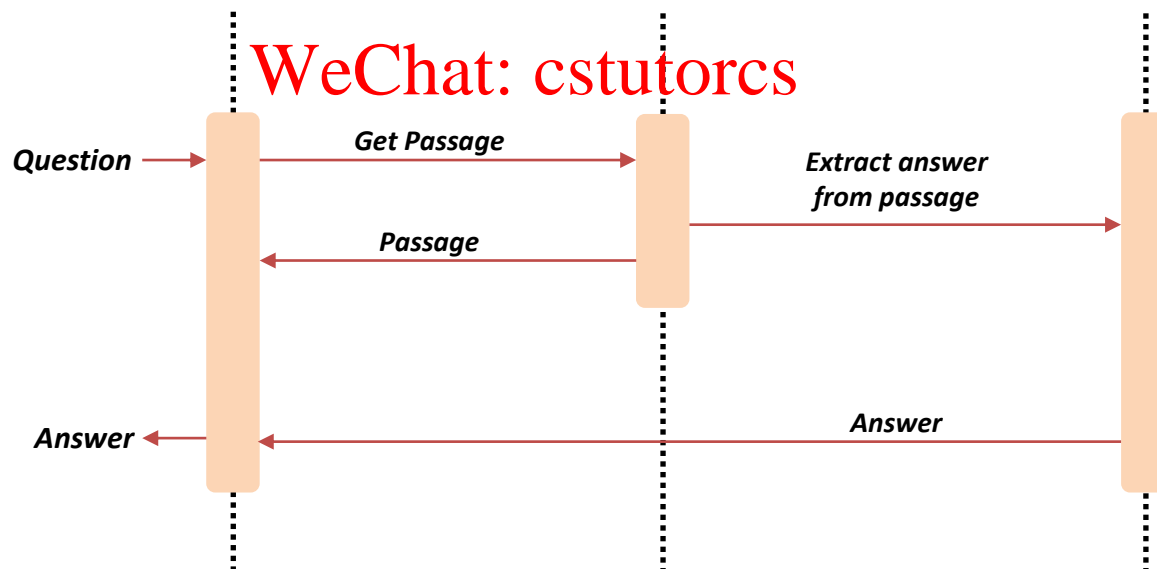- *Difficult to find the well-structured training dataset*

Research Areas in Question Answering

Research Area	Details
Knowledge-based QA (Semantic Parsing)	• Answer is a logical form, possible executed against a Knowledge Base • Context is a Knowledge Base
Information Retrieval-based QA • Answer sentence selection • Reading Comprehension	• Answer is a document, paragraph, sentence • Context is a corpus of documents or a specific document
Visual QA	• Answer is simple and factual • Context is one/multiple image(s)
Library Reference	• Answer is another question • Context is the structured knowledge available in the library and the librarians view of it.

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

Information Retrieval-based Question Answering

Information Retrieval-based Question Answering

Answering a user's question by finding short text segments, sentences, or documents on the web or collection of document

Assignment Project Exam Help

Q&A Service

**IR System
(e.g. wiki/google search)**

**Machine Reading
Comprehension**

https://tutorcs.com

WeChat: cstutorcs

Question →

Get Passage →

*Extract answer
from passage* →

← *Passage*

Answer ←   ← *Answer*

3 **Reading Comprehension**

Information Retrieval-based Question Answering

Answering a user's question by finding short text segments, sentences, or documents on the web or collection of document

Assignment Project Exam Help

- *Reading Comprehension and Answer Sentence Selection:*
  - – *Finding an answer in a paragraph or a document*
    https://tutorcs.com
 - – *Picking a suitable sentence from a corpus that can be used to answer a question*
    WeChat: cstutorcs

Reading Comprehension

To answer these questions, you need to first gather information by collecting answer-related sentences from the article.

Assignment Project Exam Help

Can we teach this to machine?

https://tutorcs.com

WeChat: cstutorcs

Yes, we can!
Machine Comprehension of Text
(Burges 2013)

THE BOAT PARADE

The boats are floating along the lakeshore. It is the summer boat parade.

There are motor boats, rowboats and sailboats.

Lisa's favorite is the yellow motor boat with the flag. The rowboat decorated with flowers is Lisa's favorite. Tony likes the purple sailboat.

The boats float by one at a time. The people on the boats waive at the crowds. The crowds cheer the boats.

The boat parade is so much fun to watch. It is the best part of the summer.

Answer the Questions:

1. Where are the boats floating?

2. What kind of boats are there?

3. What is Lisa's favorite boat?

Reading Comprehension

To answer these questions, you need to first gather information by collecting answer-related sentences from the article.

Assignment Project Exam Help

https://tutorcs.com

Can we teach this to machine?

WeChat: cstutorcs

A machine comprehends a passage of text if, for any question regarding that text that can be answered correctly by a majority of native speakers

THE BOAT PARADE

The boats are floating along the lakeshore. It is the summer boat parade.

There are motor boats, rowboats and sailboats.

Lisa's favorite is the yellow motor boat with the flag. The rowboat decorated with flowers is Lisa's favorite. Tony likes the purple sailboat.

The boats float by one at a time. The people on the boats waive at the crowds. The crowds cheer the boats.

The boat parade is so much fun to watch. It is the best part of the summer.

Answer the Questions:

1. Where are the boats floating?

2. What kind of boats are there?

3. What is Lisa's favorite boat?

Reading Comprehension

Reading Comprehension

To answer these questions, you need to first gather information by collecting answer-related sentences from the article.

Assignment Project Exam Help

https://tutorcs.com

Why do we need to teach this?

WeChat: cstutorcs

The ability to comprehend text will lead us to a better search and solve lots of NLP problems!

THE BOAT PARADE

The boats are floating along the lakeshore. It is the summer boat parade.

There are motor boats, rowboats and sailboats.

Lisa's favorite is the yellow motor boat with the flag. The rowboat decorated with flowers is Lisa's favorite. Tony likes the purple sailboat.

The boats float by one at a time. The people on the boats waive at the crowds. The crowds cheer the boats.

The boat parade is so much fun to watch. It is the best part of the summer.

Answer the Questions:

1. Where are the boats floating?

2. What kind of boats are there?

3. What is Lisa's favorite boat?

Reading Comprehension

Corpora for Reading Comprehension

Dataset	Answer Type	Domain
MCTest (Richardson et al. 2013)	Multiple choice	Children's stories
CNN/Daily Mail (Hermann et al. 2015)	Spans	News
Children's book test (Hill et al. 2016)	Multiple choice	Children's stories
SQuAD (Rajpurkar et al., 2016)	Spans	Wikipedia
MS MARCO (Nguyen et al., 2016)	Free-from text, Unanswerable	Web Search
NewsQA (Trischler et al., 2017)	Spans	News
SearchQA (Dunn et al., 2017)	Spans	Jeopardy
TriviaQA (Joshi et al., 2017)	Spans	Trivia
RACE (Lai et al., 2017)	Multiple choice	Mid/High School Exams
Narrative QA (Kocisky et al., 2018)	Free-form text	Movie Scripts, Literature
SQuAD 2.0 (Rajpurkar et al., 2018)	Spans, Unanswerable	Wikipedia

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

Reading Comprehension

TriviaQA: A Large Scale Dataset for Reading Comprehension

TriviaQA: A Large Scale Distantly Supervised Challenge Dataset for Reading Comprehension

The full dataset is coming soon. Here's a sneak peek! The evidence documents come from two domains -- Wikipedia and the web. Click on the "Evidence" button to see the document for each question.

QuestionId	Question	Answer	Web	Wikipedia
qw_3199	Miami Beach in Florida borders which ocean?	Atlantic	Evidence	Evidence
bt_1255	What was the occupation of Lovely Rita according to the song by the Beatles	Traffic Warden	Evidence	Evidence
qg_77	Who was Poopdeck Pappys most famous son?	Popeye	Evidence	Evidence
wh_1026	The Nazi regime was Germany's Third Reich; which was the first Reich?	HOLY ROMAN EMPIRE	Evidence	Evidence
bb_1342	At which English racecourse did two horses collapse and die in the parade ring due to electrocution, in February 2011?	Newbury	Evidence	Evidence
wh_2759	Which type of hat takes its name from an 1894 novel by George Du Maurier where the title character has the surname O'Ferrall ?	TRILBY	Evidence	Evidence
sfq_8522	What was the Elephant Man's real name?	Joseph Merrick	Evidence	Evidence

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

http://nlp.cs.washington.edu/triviaqa/sample.html

Reading Comprehension

TriviaQA: A Large Scale Dataset for Reading Comprehension

Our UsydNLP achieved the **No.1** in the **TriviaQA Leaderboard (Web Setting)**!

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

Wikipedia	Web

Phase description
This phase refers to the Web domain described in the paper.

Max submissions per day: 3

Max submissions total: 100

Download CSV

Results

#	User	Entries	Date of Last Entry	Team Name	full-em ▲	full-f1 ▲	verified-em ▲	verified-f1 ▲
1	mandarjoshi	10	12/02/19	Dataset Author (Oracle Result)	82.99 (1)	87.18 (1)	90.38 (1)	92.96 (1)
2	usydnlp	11	01/11/21		72.17 (2)	77.36 (2)	84.40 (2)	87.11 (2)
3	NEUKG	9	08/12/19		69.64 (3)	73.80 (3)	83.36 (3)	85.66 (4)
4	mingyan	1	02/26/18	SLQA	68.65 (4)	73.07 (5)	82.44 (5)	85.35 (5)
5	scv.back	4	06/07/18	S3R	68.21 (5)	73.26 (4)	82.57 (4)	86.05 (3)
6	dirkweissenborn	7	01/15/18		67.46 (6)	72.80 (6)	77.63 (9)	82.01 (8)
7	S3R	5	02/28/18		66.82 (7)	71.91 (7)	81.01 (6)	84.12 (6)
8	chrisc	3	09/24/17		66.37 (8)	71.32 (8)	79.97 (7)	83.70 (7)
9	Mary	5	05/26/20		66.09 (9)	69.78 (9)	79.71 (8)	81.88 (9)
10	shuohang	3	11/13/17		63.04 (10)	68.53 (10)	69.70 (10)	74.57 (10)
11	swabha	1	10/27/17	swabha_ankur_tom	53.75 (11)	58.57 (11)	63.20 (11)	66.88 (11)
12	hux444	2	07/25/17		46.65 (12)	52.89 (12)	56.96 (12)	61.48 (12)

http://nlp.cs.washington.edu/triviaqa/sample.html

Reading Comprehension

SQuAD: Stanford Question Answering Dataset

Victoria_(Australia)
The Stanford Question Answering Dataset

The economy of Victoria is highly diversified: service sectors including financial and property services, health, education, wholesale, retail, hospitality and manufacturing constitute the majority of employment. Victoria's total gross state product (GSP) is ranked second in Australia, although Victoria is ranked fourth in terms of GSP per capita because of its limited mining activity. Culturally, Melbourne is home to a number of museums, art galleries and theatres and is also described as the "sporting capital of Australia". The Melbourne Cricket Ground is the largest stadium in Australia, and the host of the 1956 Summer Olympics and the 2006 Commonwealth Games. The ground is also considered the "spiritual home" of Australian cricket and Australian rules football, and hosts the grand final of the Australian Football League (AFL) each year, usually drawing crowds of over 95,000 people. Victoria includes eight public universities, with the oldest, the University of Melbourne, having been founded in 1853.

What kind of economy does Victoria have?
Ground Truth Answers: diversified   highly diversified   highly diversified
Prediction: highly diversified

Where according to gross state product does Victoria rank in Australia?
Ground Truth Answers: second   second   second
Prediction: second

At what rank does GPS per capita set Victoria?
Ground Truth Answers: fourth   fourth   fourth
Prediction: fourth

What city in Victoria is called the sporting capital of Australia?
Ground Truth Answers:
Melbourne   Melbourne   Melbourne
Prediction: Melbourne

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

https://rajpurkar.github.io/SQuAD-explorer/

Reading Comprehension

A Generic Neural Model for Reading Comprehension

Step1: For both documents and questions, convert words to word vectors

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

A   partly   …   Perito   Moreno   …   How   are   …

Document (D)

A partly submerged glacier cave on Perito Moreno Glacier. The ice facade is approximately 60 m high. Ice formations in the Titlis glacier cave. A glacier cave is a cave formed within the ice of a glacier. Glacier caves are often called ice caves, but the latter term is properly used to describe bedrock caves that contain year-round ice

Question (Q)

How are glacier caves formed?

A Generic Neural Model for Reading Comprehension

Step2: *Encode* *context (documents)* *and* *question* *with sequence models*

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

Sequence Model

Sequence Model

A   partly   …   Perito   Moreno   …

How   are   …

Document (D)

A partly submerged glacier cave on Perito Moreno Glacier. The ice facade is approximately 60 m high. Ice formations in the Titlis glacier cave. A glacier cave is a cave formed within the ice of a glacier. Glacier caves are often called ice caves, but the latter term is properly used to describe bedrock caves that contain year-round ice

Question (Q)

How are glacier caves formed?

Reading Comprehension

A Generic Neural Model for Reading Comprehension

Step3: *Combine* *context (documents)* *and* *question* *with an attention*

Attention

Sequence Model

Sequence Model

A   partly   ...   Perito   Moreno   ...   How   are   ...

Document (D)

A partly submerged glacier cave on Perito Moreno Glacier. The ice facade is approximately 60 m high. Ice formations in the Titlis glacier cave. A glacier cave is a cave formed within the ice of a glacier. Glacier caves are often called ice caves, but the latter term is properly used to describe bedrock caves that contain year-round ice
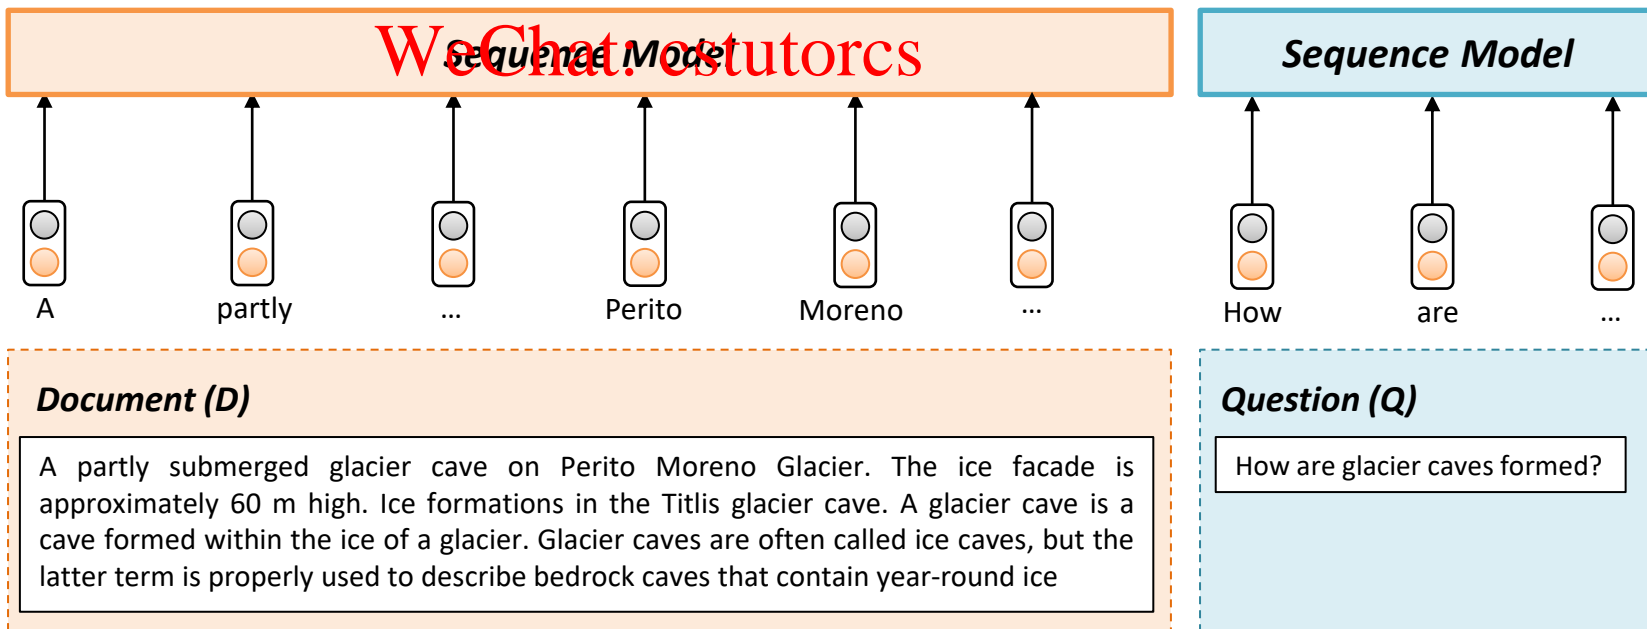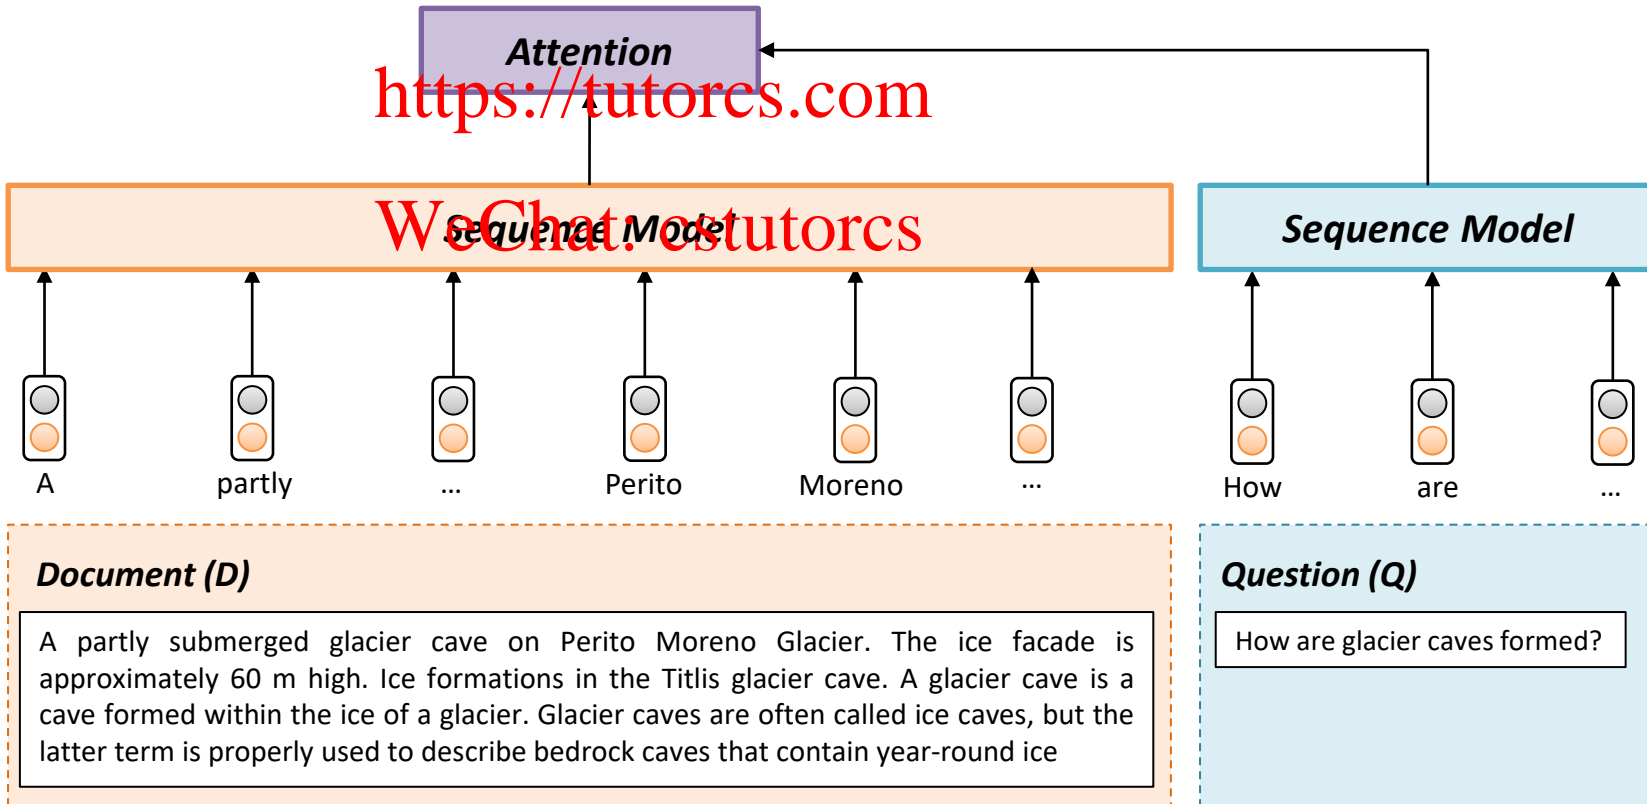
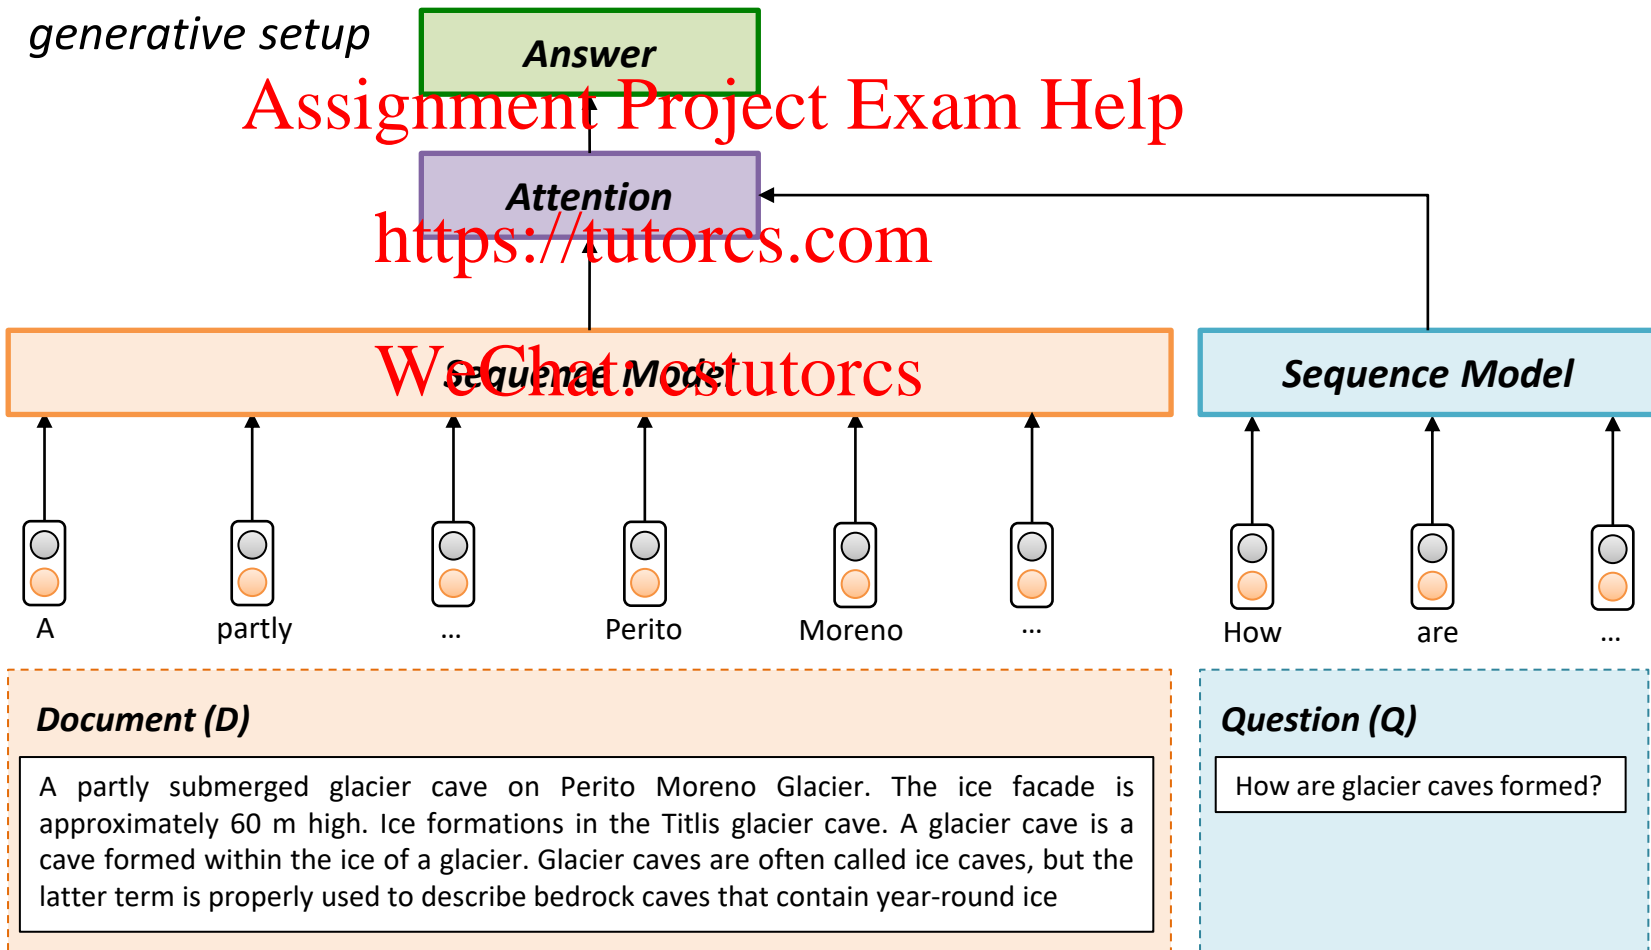**Question (Q)**

How are glacier caves formed?
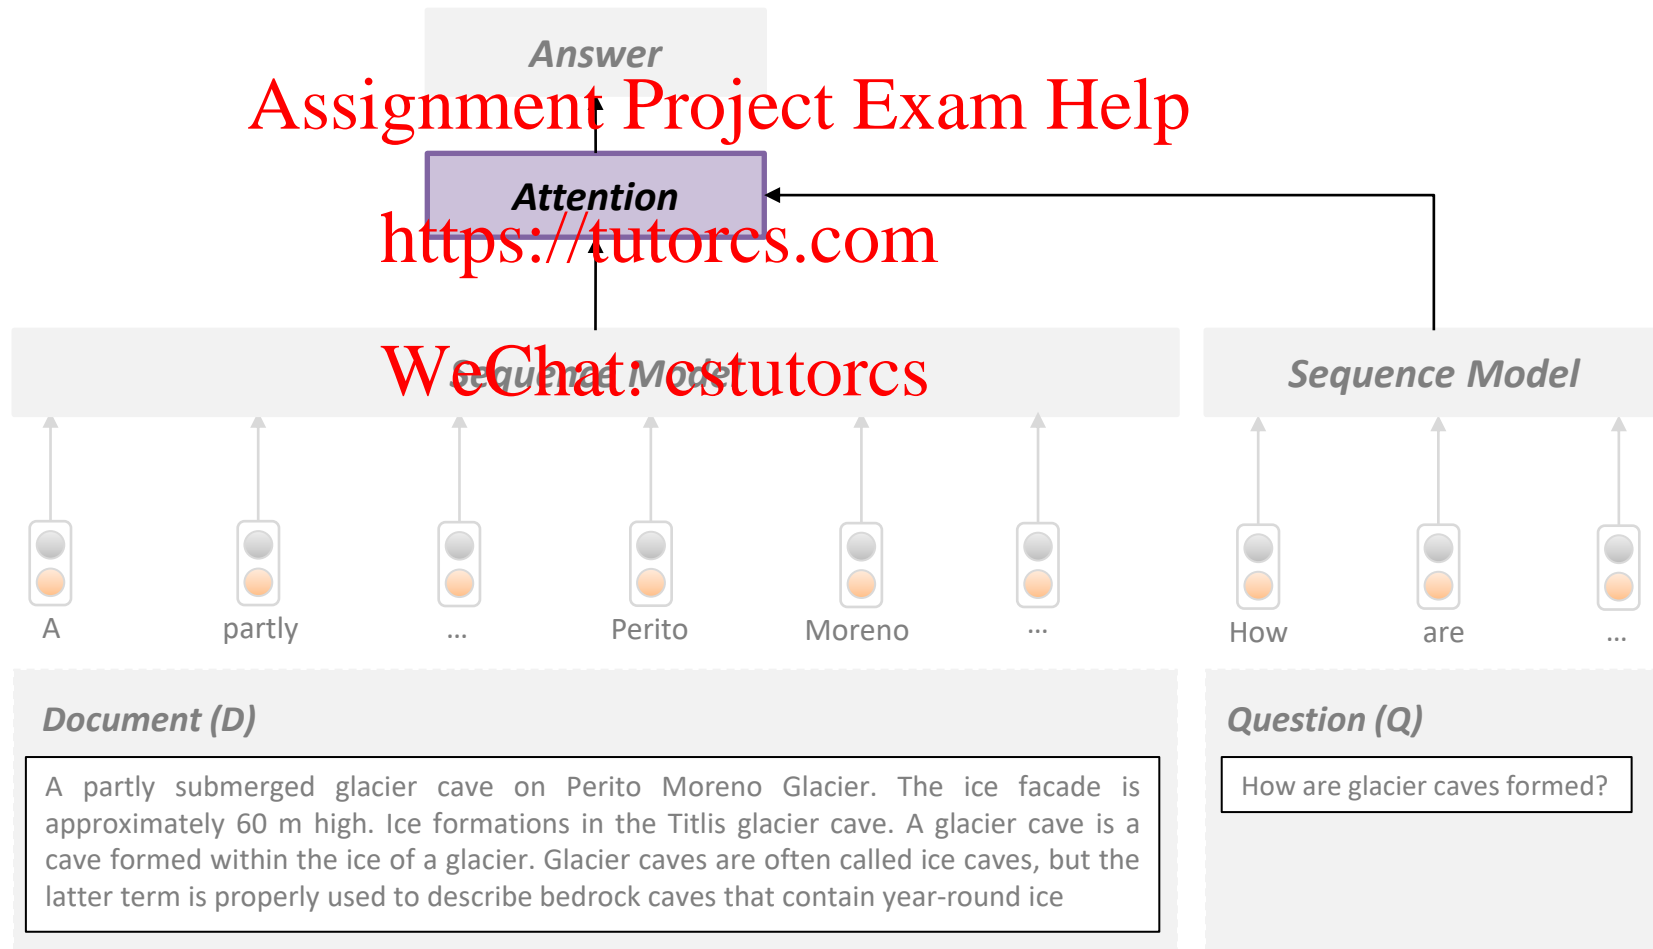
Reading Comprehension

A Generic Neural Model for Reading Comprehension

Step4: *Select answer from attention map by using a classifier or with generative setup*

Answer

Assignment Project Exam Help

Attention

https://tutorcs.com

Sequence Model

WeChat: cstutorcs

Sequence Model

A   partly   …   Perito   Moreno   …   How   are   …

Document (D)

A partly submerged glacier cave on Perito Moreno Glacier. The ice facade is approximately 60 m high. Ice formations in the Titlis glacier cave. A glacier cave is a cave formed within the ice of a glacier. Glacier caves are often called ice caves, but the latter term is properly used to describe bedrock caves that contain year-round ice

Question (Q)

How are glacier caves formed?

A Generic Neural Model for Reading Comprehension

*What is the **Attention**? Why we need this?*

Answer

Assignment Project Exam Help

Attention

https://tutorcs.com

Sequence Model   *Sequence Model*

WeChat: cstutorcs

A   partly   …   Perito   Moreno   …   How   are   …

Document (D)

A partly submerged glacier cave on Perito Moreno Glacier. The ice facade is approximately 60 m high. Ice formations in the Titlis glacier cave. A glacier cave is a cave formed within the ice of a glacier. Glacier caves are often called ice caves, but the latter term is properly used to describe bedrock caves that contain year-round ice

Question (Q)

How are glacier caves formed?

Seq2Seq Model: Recap

[Encoder]

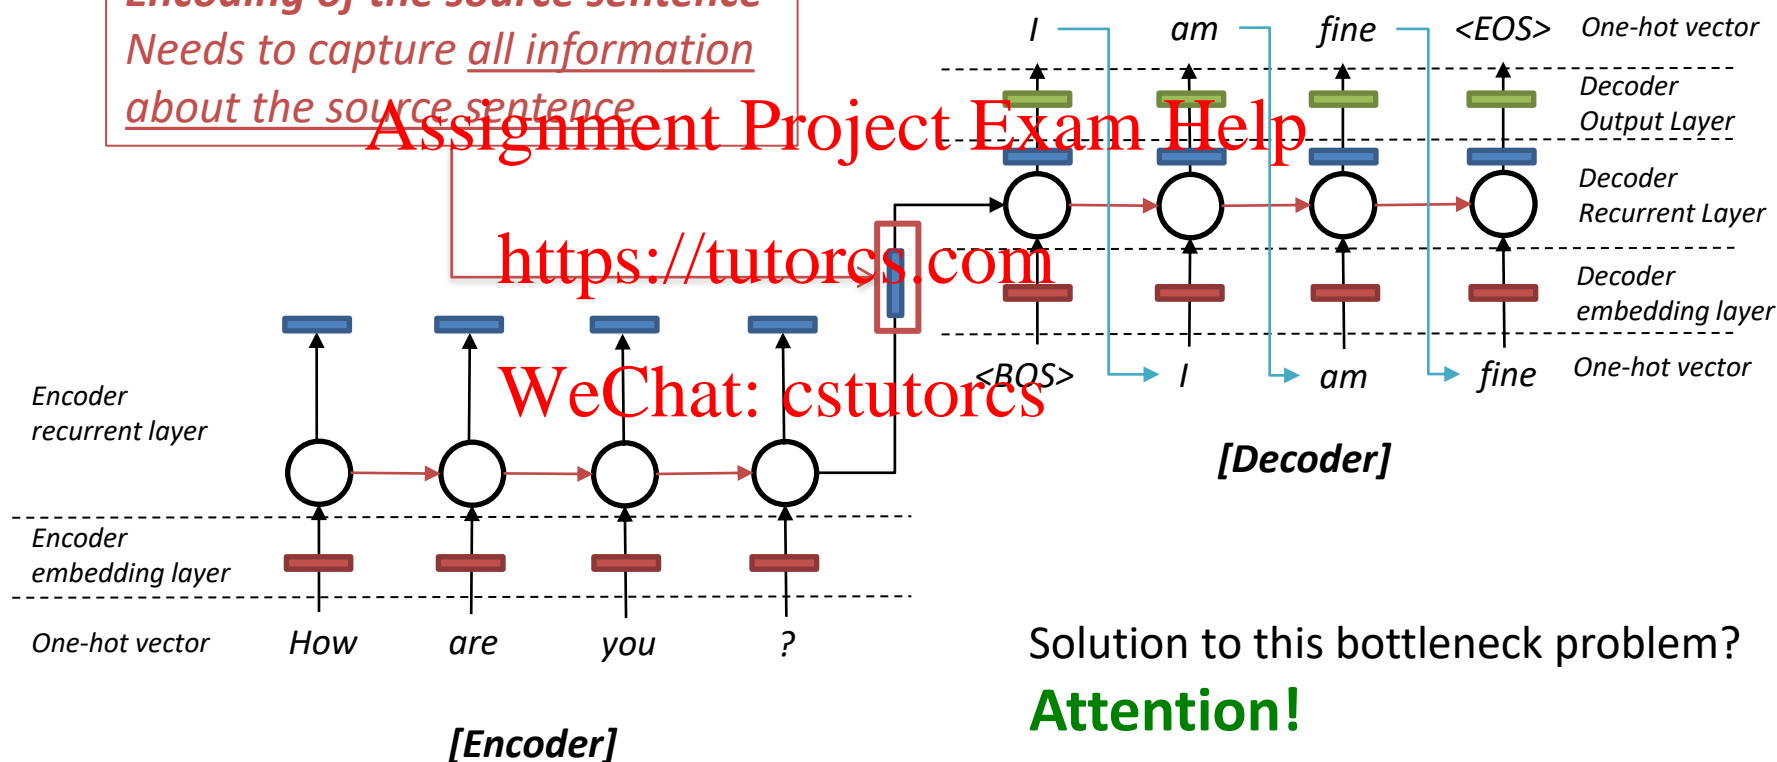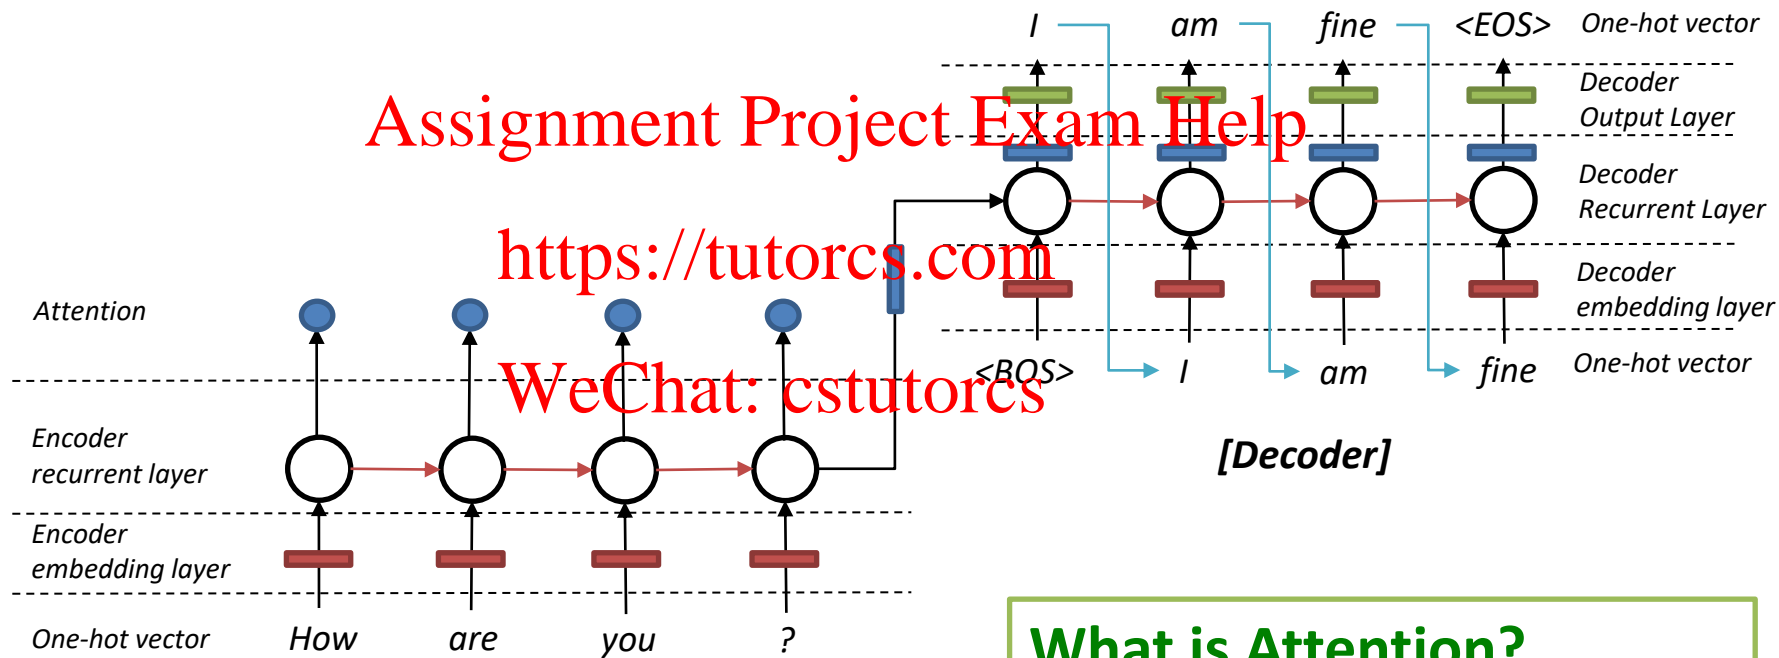*[Decoder]*

Seq2Seq Model: the bottleneck problem

Encoding of the source sentence
Needs to capture all information
about the source sentence

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

I   *am*   *fine*   *<EOS>*   One-hot vector

Decoder
Output Layer

Decoder
Recurrent Layer

Decoder
embedding layer

<BOS>   *I*   *am*   *fine*   One-hot vector

[Decoder]

Encoder
recurrent layer

Encoder
embedding layer

One-hot vector   *How*   *are*   *you*   *?*

[Encoder]

+RNN drawback!
Vanishing Gradient

Solution to this bottleneck problem?
Attention!

Attention

Seq2Seq with Attention

[Decoder]

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

What is Attention?
*On **each step of the decoder**, use **direct connection to the encoder** to **focus on a particular part** of the input sequence*

**Seq2Seq with Attention**

Assignment Project Exam Help

dot product https://tutorcs.com

WeChat: cstutorcs

Attention scores

Encoder recurrent layer

Decoder Recurrent Layer

Encoder embedding layer

Decoder embedding layer

One-hot vector   How   are   you   ?   <BOS>   One-hot vector

4 **Attention**

Seq2Seq with Attention

Assignment Project Exam Help

dot product https://tutorcs.com

Attention scores

WeChat: cstutorcs

Encoder recurrent layer

Decoder Recurrent Layer

Encoder embedding layer

Decoder embedding layer

One-hot vector How are you ? <BOS> *One-hot vector*

Seq2Seq with Attention

Take softmax to turn the scores
into a probability distribution

Assignment Project Exam Help

softmax

Attention
distribution

https://tutorcs.com

Attention
scores

WeChat: cstutorcs

Encoder
recurrent layer

Decoder
Recurrent Layer

Encoder
embedding layer

Decoder
embedding layer

One-hot vector   How   are   you   ?   <BOS>   One-hot vector

Attention

Seq2Seq with Attention

On this decoder step, focusing on the encoder hidden state

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

Attention distribution

softmax

Attention scores

Encoder recurrent layer

Decoder Recurrent Layer

Encoder embedding layer

Decoder embedding layer

One-hot vector   How   are   you   ?   <BOS>   *One-hot vector*

Attention
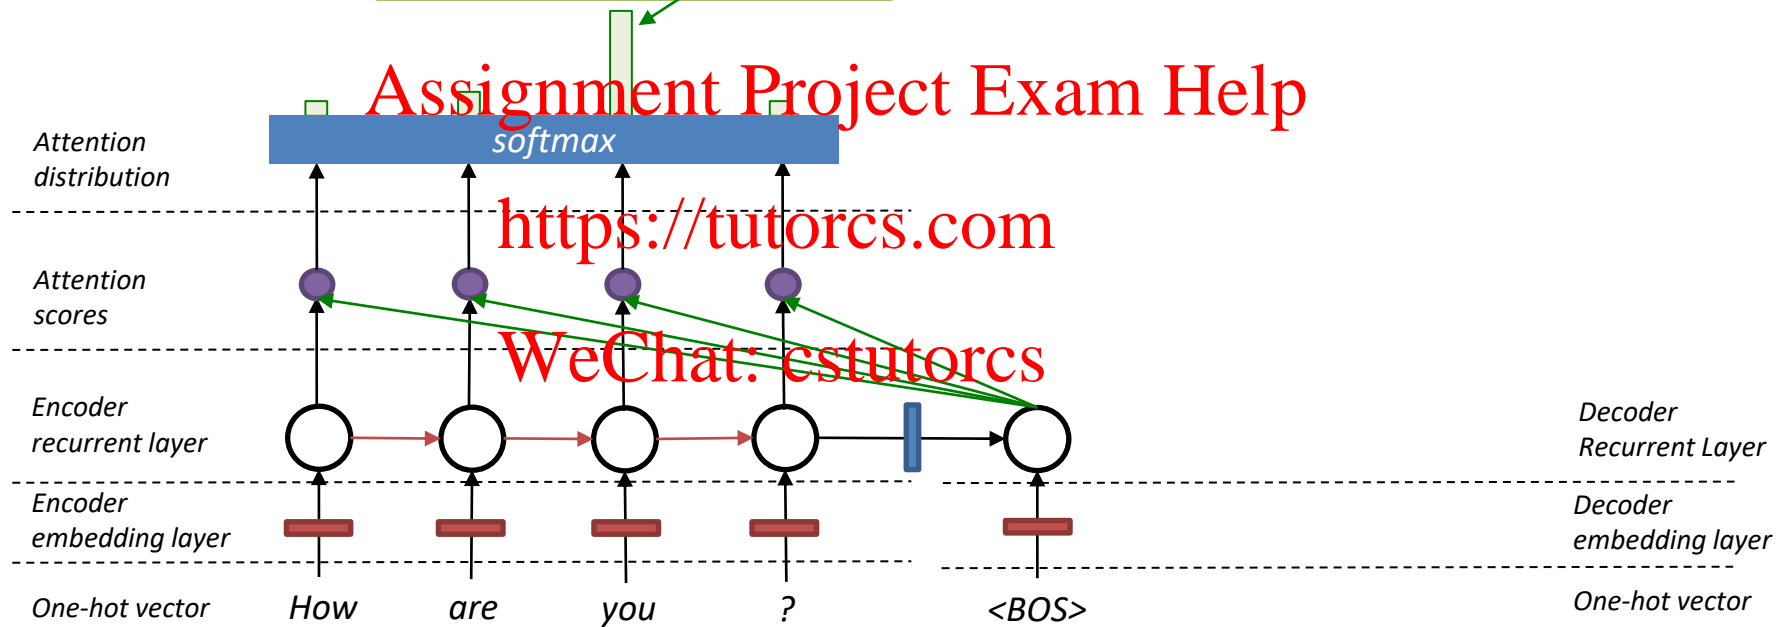
Seq2Seq with Attention

Attention output

The attention output mostly contains information from the hidden states that received high attention.

Calculate a **weighted sum** of the *encoder hidden states* with *the attention distribution*

Assignment Project Exam Help

Attention distribution

softmax

https://tutorcs.com

Attention scores

WeChat: cstutorcs

Encoder recurrent layer

Decoder Recurrent Layer

Encoder embedding layer

Decoder embedding layer

One-hot vector   How   are   you   ?   <BOS>   *One-hot vector*

Seq2Seq with Attention

Attention output

Concatenate attention output with decoder hidden state

Attention distribution

softmax

Attention scores

Encoder recurrent layer

Encoder embedding layer

One-hot vector

How   are   you   ?   <BOS>

Decoder Recurrent Layer

Decoder embedding layer

One-hot vector

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

Attention

Seq2Seq with Attention

am

Attention
output

Assignment Project Exam Help

Attention
distribution

softmax

https://tutorcs.com

Attention
scores

WeChat: cstutorcs

Encoder
recurrent layer

*Decoder
Recurrent Layer*

Encoder
embedding layer

*Decoder
embedding layer*

One-hot vector *How*   *are*   *you*   *?*   *<BOS>*   *I*   *One-hot vector*

Seq2Seq with Attention

fine

Attention output

Assignment Project Exam Help

Attention distribution

softmax

https://tutorcs.com

Attention scores

WeChat: cstutorcs

Encoder recurrent layer

Decoder Recurrent Layer

Encoder embedding layer

Decoder embedding layer

One-hot vector   How   are   you   ?   <BOS>   I   am   *One-hot vector*

Attention

THE UNIVERSITY OF SYDNEY

Seq2Seq with Attention

<EOS>

Attention output

Attention distribution

softmax

Assignment Project Exam Help

https://tutorcs.com

Attention scores

WeChat: cstutorcs

Encoder recurrent layer

Decoder Recurrent Layer

Encoder embedding layer

Decoder embedding layer

One-hot vector

| How | are | you | ? | | <BOS> | I | am | fine |

One-hot vector

Attention

THE UNIVERSITY OF SYDNEY

Seq2Seq with Attention: Machine Translation

!

Attention output

Assignment Project Exam Help

Attention distribution

softmax

https://tutorcs.com

Attention scores

WeChat: cstutorcs

Encoder recurrent layer

Decoder Recurrent Layer

Encoder embedding layer

Decoder embedding layer

One-hot vector

I   love   you   !   <BOS>   我   爱   你

One-hot vector

Seq2Seq with Attention (Equations)

- *Encoder hidden states:* $h_1, \ldots, h_N \in \mathbb{R}^h$

- *Decoder hidden state:* $s_t \in \mathbb{R}^h$ *(on timestep t)*

1. *Attention score* $\boldsymbol{e}^t$ *:* $\boldsymbol{e}^t = [\boldsymbol{s}_t^T \boldsymbol{h}_1, \ldots, \boldsymbol{s}_t^T \boldsymbol{h}_N] \in \mathbb{R}^N$ *(for timestep t)*

2. *Use softmax to get the attention distribution,* α^t *(for timestep t)*
   (this is a probability distribution and sums to 1)

$$\alpha^t = \text{softmax}(\boldsymbol{e}^t) \in \mathbb{R}^N$$

3. *Attention Output: Use* α^t *to take a weighted sum of the encoder hidden states*

$$\boldsymbol{a}_t = \sum_{i=1}^{N} \alpha_i^t \boldsymbol{h}_i \in \mathbb{R}^h$$

4. *Then, concatenate the attention output* $\boldsymbol{a}_t$ *with the decoder hidden state* s_t *and proceed as in the non-attention seq2seq model*

$$[\boldsymbol{a}_t; \boldsymbol{s}_t] \in \mathbb{R}^{2h}$$

Attention

Why we use Attention? The benefit!

Improve performance

- *Allow decoder to focus on certain parts of the source*

Assignment Project Exam Help

Solving the bottleneck problem

- *Allow decoder to directly look at the source (input)*

https://tutorcs.com

WeChat: cstutorcs

Reducing vanishing gradient problem

- *Provide shortcut to faraway states*

Providing some interpretability

- *Inspect attention distribution, and show what the decoder was focusing on*

Attention

Attention is now a general component in Deep Learning NLP

Attention is great way to improve the sequence to sequence model.

You can use attention in many architecture (not just seq2seq) and many NLP tasks (e.g. summarization / text classification / translation).

More general definition of attention:

Given a set of vector values, and a vector query, attention is a technique to compute a weighted sum of the values, dependent on the query.

*For example, in the seq2seq + attention model, each decoder hidden state (**query**) attends to all the encoder hidden states (**values**).*
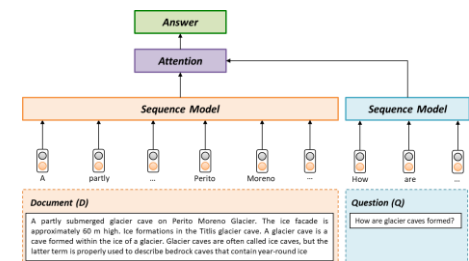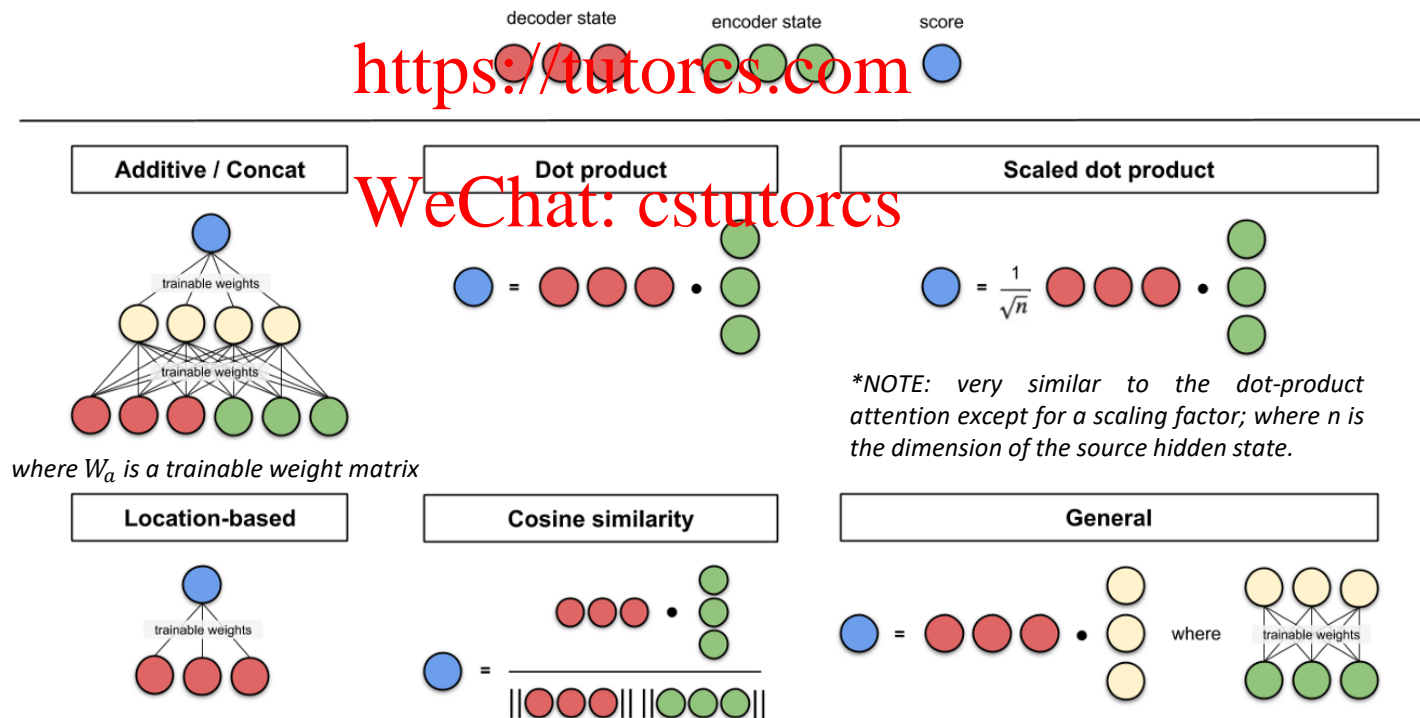
Attention variants

There are several ways to compute attention score.

- Encoder hidden states: $h_1, \ldots, h_N \in \mathbb{R}^h$
- Decoder hidden state: $s_t \in \mathbb{R}^h$ (on timestep t)

Attention Name	Attention score function	Reference
Content-base	$score(s_t, h_i) = cosine[s_t, h_i]$	Graves 2014
Dot-product	$score(s_t, h_i) = s_t^\top h_i$	Luong 2015
Scaled Dot-product	$score(s_t, h_i) = \frac{s_t^\top h_i}{\sqrt{n}}$ *NOTE: very similar to the dot-product attention except for a scaling factor; where n is the dimension of the source hidden state.	Vaswani 2017
Additive	$score(s_t, h_i) = v_a^\top \tanh(W_a[s_t; h_i])$	Vaswani 2017
General	$score(s_t, h_i) = s_t^\top W_a h_i$ *NOTE: where W_a is a trainable weight matrix in the attention layer.	Luong 2015
Location-based	$a_{t,i} = softmax(W_a s_t)$ *Note: This simplifies the softmax alignment to only depend on the target position.	Luong 2015

The papers (Luong 2015 and Vaswani 2017) can be found in the canvas content page

Attention

Attention variants

There are several ways to compute attention score.

- *Encoder hidden states:* $h_1, \ldots, h_N \in \mathbb{R}^h$
- *Decoder hidden state:* $s_t \in \mathbb{R}^h$ *(on timestep t)*

decoder state   encoder state   score

| Additive / Concat | Dot product | Scaled dot product |

$\frac{1}{\sqrt{n}}$

NOTE: very similar to the dot-product attention except for a scaling factor; where n is the dimension of the source hidden state.

trainable weights

trainable weights

where W_a is a trainable weight matrix

| Location-based | Cosine similarity | General |

trainable weights

where   trainable weights

Note: This simplifies the softmax alignment to only depend on the target position.

Categories of Attention Mechanism

A summary of broader categories of attention mechanisms

Name	Definition	Citation
Global or Local	• Global: Attending to the entire input state space. • Local: Attending to the part of input state space (i.e. a patch of the input image.)	**Luong 2015**
Self-Attention	Relating different positions of the same input sequence. Theoretically the self-attention can adopt any attention score functions, but just replace the target sequence with the same input sequence.	**Cheng 2016**

*The papers (Luong 2015 and Cheng 2016) can be found in the canvas content page

Categories of Attention Mechanism (1)

Global/Local Attention

- *Global: Attending to the entire input state space.*

- *Local: Attending to the part of input state space*

Global Attention Model

Local Attention Model
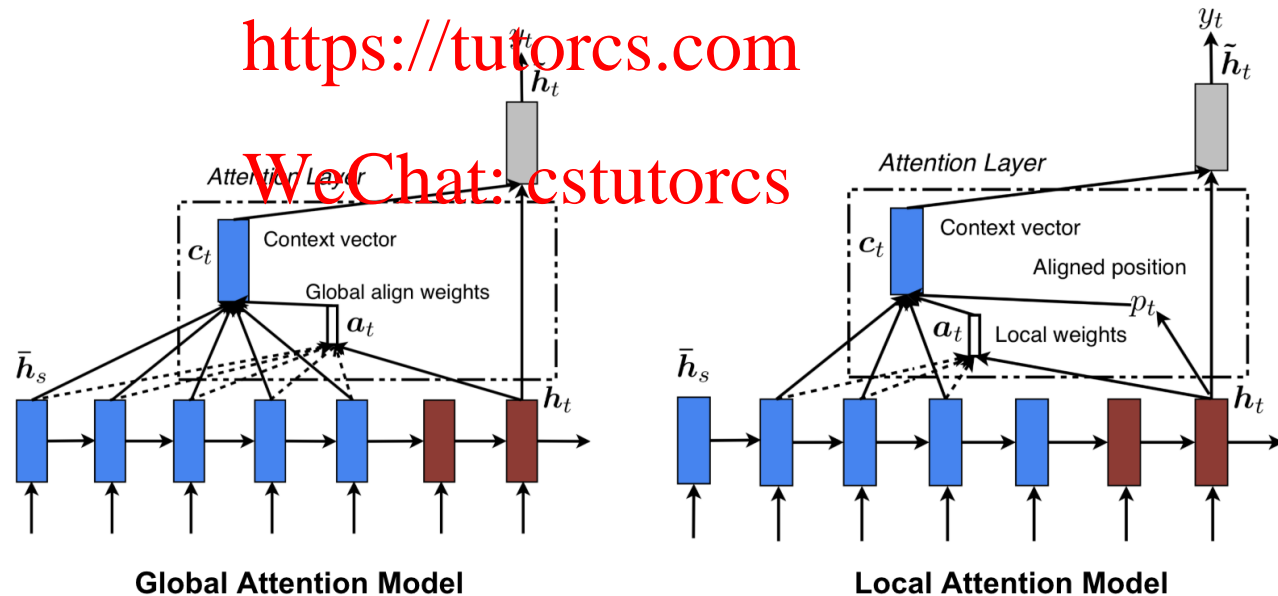
Categories of Attention Mechanism (2)

Self-Attention

The long short-term memory network (Cheng et al., 2016) paper used self-attention to do machine reading. In the example below, the self-attention mechanism enables us to learn the correlation between the current words and the previous part of the sentence.
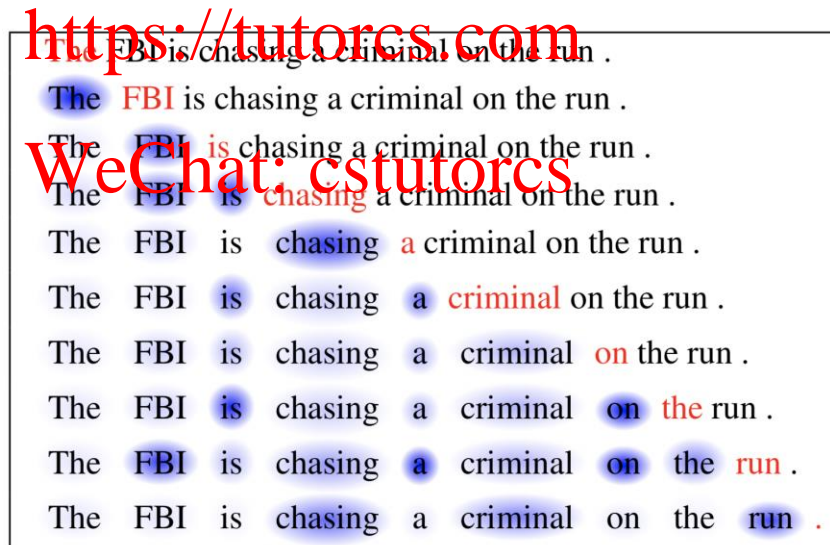
The current word is in red and the size of the blue shade indicates the activation level.

Reading Comprehension with Attention

A Generic Neural Model for Reading Comprehension

Answer

Attention

Sequence Model

Sequence Model

A   partly   …   Perito   Moreno   …

How   are   …

Document (D)

A partly submerged glacier cave on Perito Moreno Glacier. The ice facade is approximately 60 m high. Ice formations in the Titlis glacier cave. A glacier cave is a cave formed within the ice of a glacier. Glacier caves are often called ice caves, but the latter term is properly used to describe bedrock caves that contain year-round ice

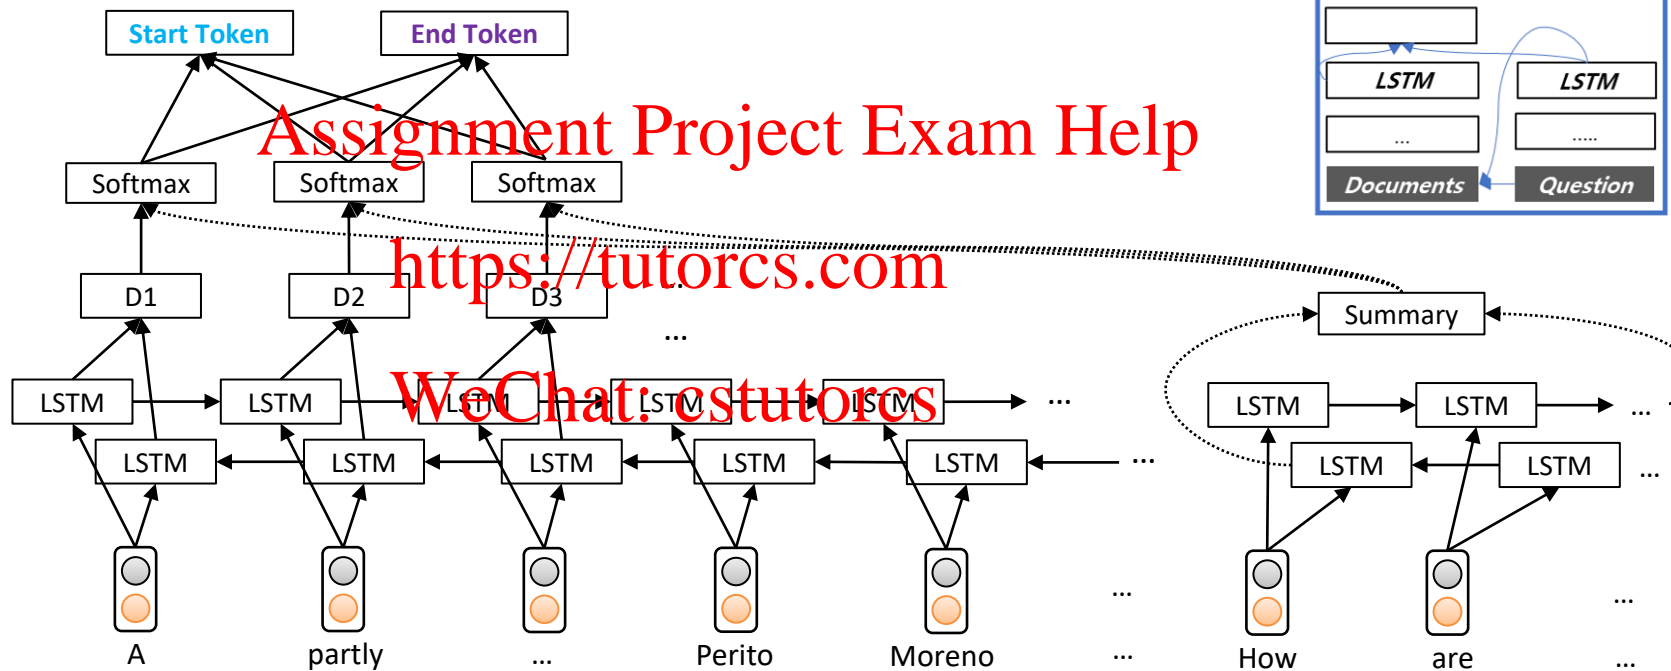**Question (Q)**

How are glacier caves formed?

Reading Comprehension with Attention

Bi-LSTM for Reading Comprehension with Attention

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

Document (D)

A partly submerged glacier cave on Perito Moreno Glacier. The ice facade is approximately 60 m high. Ice formations in the Titlis glacier cave. A glacier cave is a cave formed within the ice of a glacier. Glacier caves are often called ice caves, but the latter term is properly used to describe bedrock caves that contain year-round ice

Question (Q)

How are glacier caves formed?

Reading Comprehension with Attention

Bi-LSTM for Reading Comprehension with Attention

Predict end token?

Predict end token?

Predict end token?

Predict start token?

Predict start token?

Predict start token?

Where to put Attention?

LSTM		*LSTM*
...		......
Documents		**Question**

Softmax

Softmax

Softmax

Attention

D1

D2

D3

...

Summary

| LSTM | LSTM | LSTM | LSTM | LSTM | ... | LSTM | LSTM | ... |
| LSTM | LSTM | LSTM | LSTM | LSTM | ... | LSTM | LSTM | ... |

A   partly   ...   Perito   Moreno   ...   How   are   ...

Document (D)

A partly submerged glacier cave on Perito Moreno Glacier. The ice facade is approximately 60 m high. Ice formations in the Titlis glacier cave. A glacier cave is a cave formed within the ice of a glacier. Glacier caves are often called ice caves, but the latter term is properly used to describe bedrock caves that contain year-round ice

Question (Q)

How are glacier caves formed?

THE UNIVERSITY OF SYDNEY

Bi-LSTM for Reading Comprehension with Attention

Where to put Attention?

LSTM	LSTM
...	......
Documents	Question

Start Token

End Token

Assignment Project Exam Help

Softmax   Softmax   Softmax

https://tutorcs.com

D1   D2   D3

...

Summary

LSTM   LSTM   WeChat: LSTM cstutorcs   ...   LSTM   LSTM   ...

LSTM   LSTM   LSTM   LSTM   LSTM   ...   LSTM   LSTM   ...

A   partly   ...   Perito   Moreno   ...   How   are   ...

Document (D)

A partly submerged glacier cave on Perito Moreno Glacier. The ice facade is approximately 60 m high. Ice formations in the Titlis glacier cave. A glacier cave is a cave formed within the ice of a glacier. Glacier caves are often called ice caves, but the latter term is properly used to describe bedrock caves that contain year-round ice

Question (Q)

How are glacier caves formed?

Reading Comprehension with Attention

Bi-Directional Attention Flow (Bi-DAF)

Bi-Directional Attention Flow for Machine Comprehension (Seo et al. 2017)
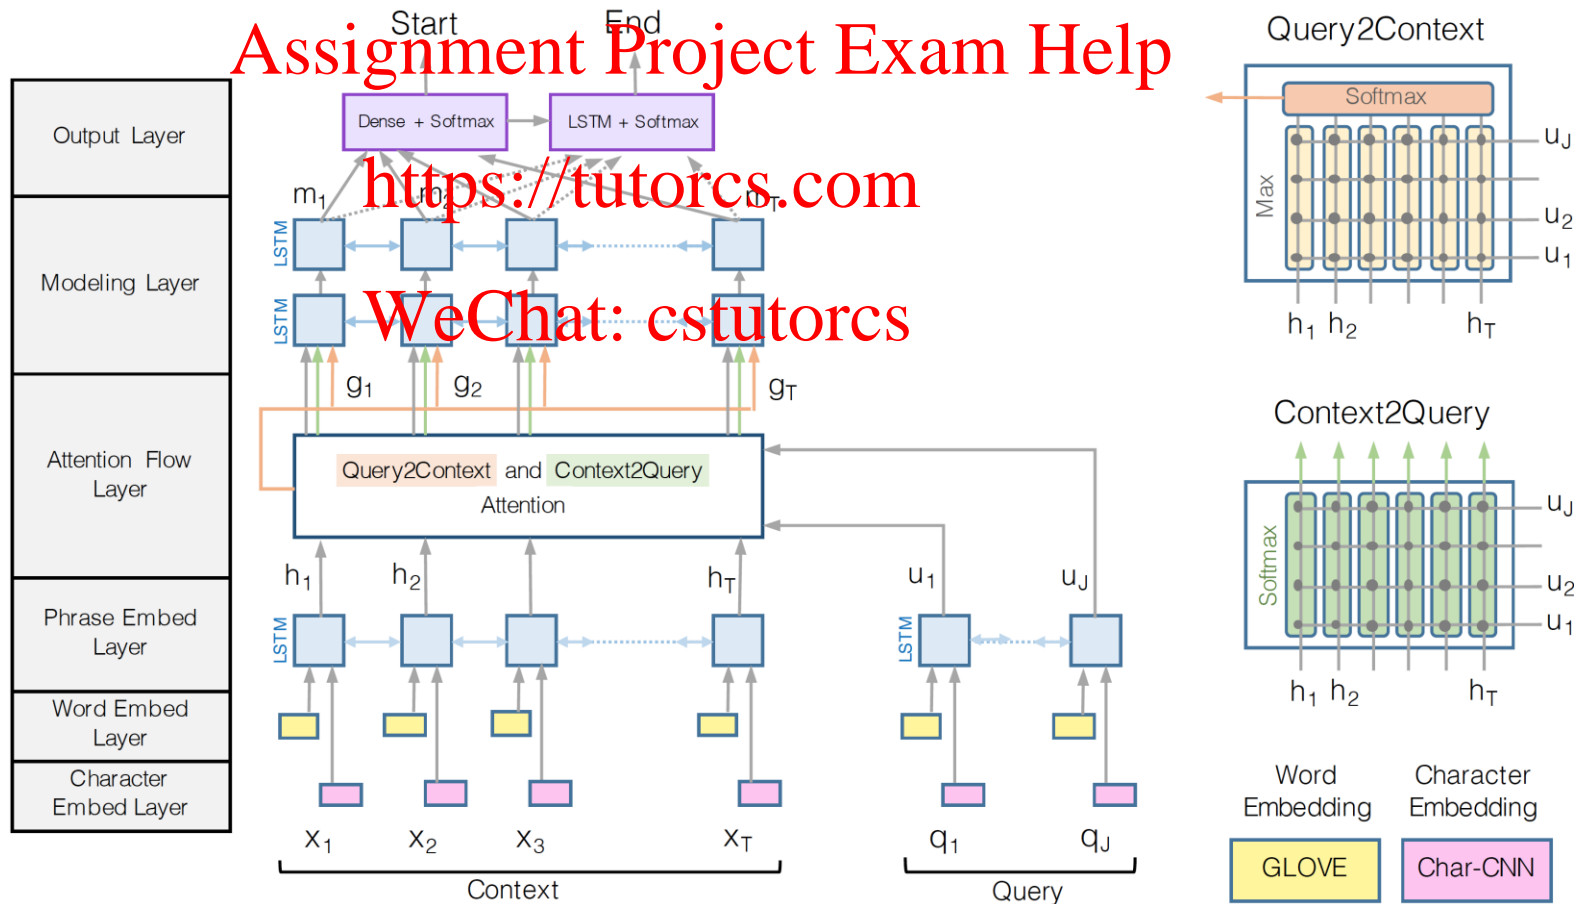
Reading Comprehension with Attention

Bi-Directional Attention Flow (Bi-DAF)

Attention Flow layer is the core idea!

- *Variants and improvements to the Bi-DAF architecture over the years*

Assignment Project Exam Help

Attention should flow both ways:

1) *the context → the question (C2Q)* https://tutorcs.com

2) *the question → the context (Q2C)*

WeChat: cstutorcs
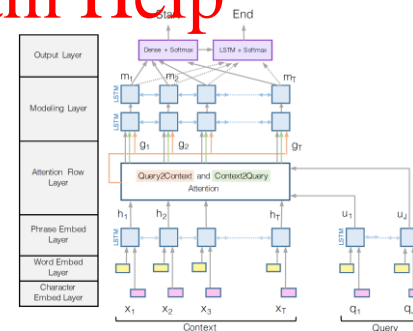
*Both attentions are derived from a **shared similarity matrix** between the **context (H) and the query (U),** where S_{tj} indicates the similarity between t-th context word and j-th query word*

$$\mathbf{S}_{tj} = \alpha(\mathbf{H}_{:t}, \mathbf{U}_{:j}) \in \mathbb{R}$$

THE UNIVERSITY OF SYDNEY

Bi-Directional Attention Flow (Bi-DAF)

Attention Flow layer is the core idea!

- *Variants and improvements to the Bi-DAF architecture over the years*

Assignment Project Exam Help

Attention should flow both ways:

1) *the context → the question (C2Q)* https://tutorcs.com

2) *the question → the context (Q2C)*

WeChat: cstutorcs

1. Context-to-Question (C2Q) attention:

- *which query words are most relevant to each context word*

$$\alpha^i = \text{softmax}(\boldsymbol{S}_{i,:}) \in \mathbb{R}^M \quad \forall i \in \{1, \dots, N\}$$

$$\boldsymbol{a}_i = \sum_{j=1}^{M} \alpha_j^i \boldsymbol{q}_j \in \mathbb{R}^{2h} \quad \forall i \in \{1, \dots, N\}$$

Context2Query

Softmax

u_J
u_2
u_1

h_1 h_2 h_T

Reading Comprehension with Attention

THE UNIVERSITY OF SYDNEY

Bi-Directional Attention Flow (Bi-DAF)

Attention Flow layer is the core idea!

- *Variants and improvements to the Bi-DAF architecture over the years*

Assignment Project Exam Help

Attention should flow both ways:

1) *the context → the question (C2Q)*
https://tutorcs.com

2) *the question → the context (Q2C)*

WeChat: cstutorcs

2. Question-to-Context (Q2C) attention:

- *the weighted sum of the most important words in the context with respect to the query – slight asymmetry through max*

$$m_i = \max_j S_{ij} \in \mathbb{R} \quad \forall i \in \{1, \ldots, N\}$$

$$\beta = \text{softmax}(\boldsymbol{m}) \in \mathbb{R}^N$$

$$\boldsymbol{c}' = \sum_{i=1}^{N} \beta_i \boldsymbol{c}_i \in \mathbb{R}^{2h}$$

Query2Context

Reading Comprehension with Attention

Bi-Directional Attention Flow (Bi-DAF)

A "modelling" layer:

- *Another deep (2-layer) Bi-LSTM over the passage*

Assignment Project Exam Help

And answer span selection is more complex:

https://tutorcs.com

- ***Start****: Pass output of BiDAF and modelling layer concatenated to a dense FF layer and then a softmax*
- ***End****: Put output of modelling layer M through another BiLSTM to give* WeChat: cstutorcs
  M_2 and then concatenate with BiDAF layer and again put through dense FF layer and a softmax
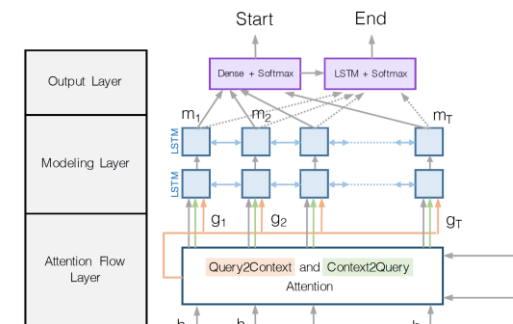
Dynamic Coattention Networks for Question Answering (Xiong 2017)

Coattention provides a two-way attention between the context and the question.

Start Token   End Token

Dynamic pointer Decoder

Co-Attention encoder

Document (D)

A partly submerged glacier cave on Perito Moreno Glacier. The ice facade is approximately 60 m high. Ice formations in the Titlis glacier cave. A glacier cave is a cave formed within the ice of a glacier. Glacier caves are often called ice caves, but the latter term is properly used to describe bedrock caves that contain year-round ice

Question (Q)

How are glacier caves formed?

Reading Comprehension with Attention

Dynamic Coattention Networks for Question Answering (Xiong 2017)

- *Coattention layer again provides a two-way attention between the context and the question*

- *Coattention involves a second-level attention computation:*
  - *attending over representations that are themselves attention outputs*
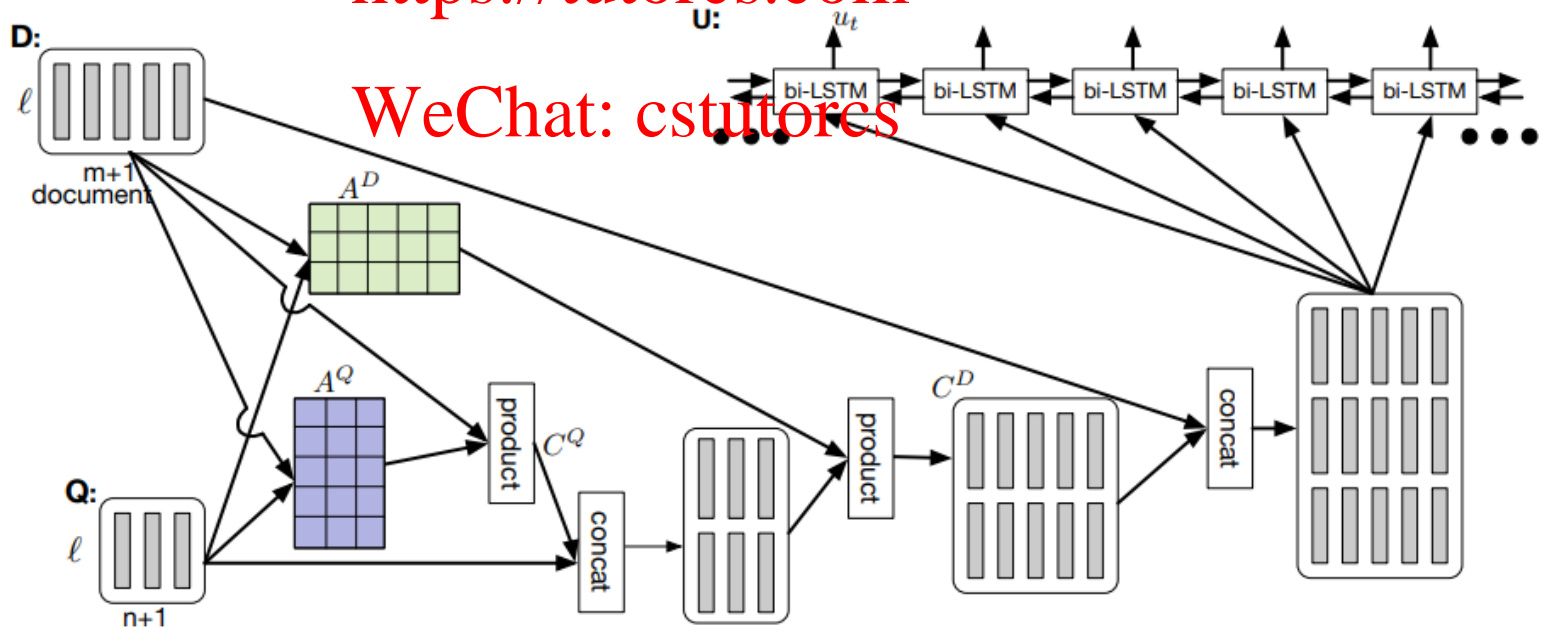
/ More…?

More Advanced Architecture? Preview for following weeks

The transformer, based solely on attention mechanisms.
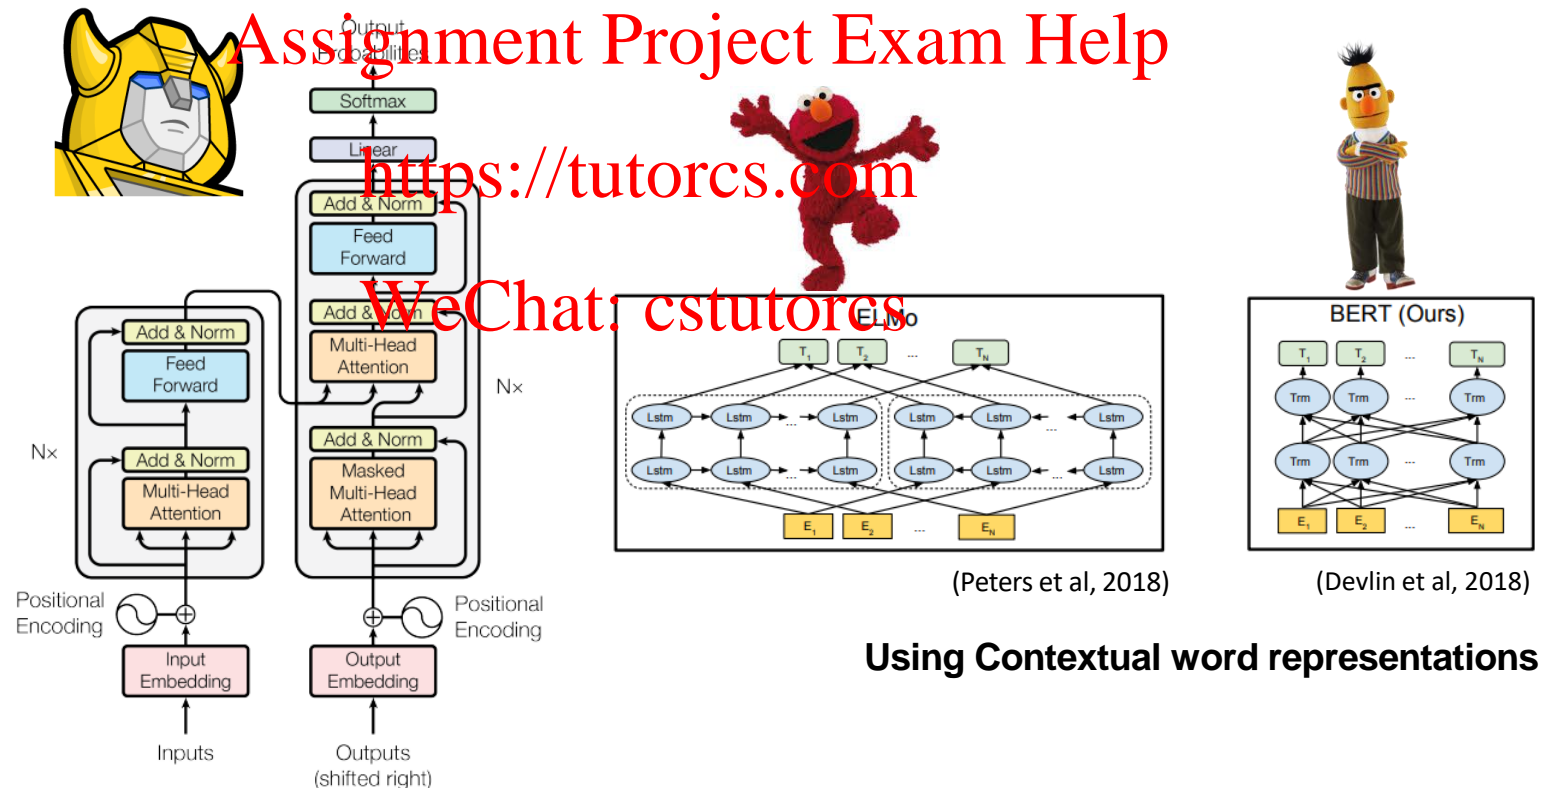
Figure 1: The Transformer - model architecture.

(Vaswani et al, 2017)

ELMo

(Peters et al, 2018)

BERT (Ours)

(Devlin et al, 2018)

Using Contextual word representations

Research Areas in Question Answering

Research Area	Details
Knowledge-based QA (Semantic Parsing)	• Answer is a logical form, possible executed against a Knowledge Base • Context is a Knowledge Base
Information Retrieval-based QA • Answer sentence selection • Reading Comprehension	• Answer is a document, paragraph, sentence • Context is a corpus of documents or a specific document
Visual QA	• **Answer is simple and factual** • **Context is one/multiple image(s)**
Library Reference	• Answer is another question • Context is the structured knowledge available in the library and the librarians view of it.

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

Visual Question Answering

Textual Question Answering: Recap

Answer questions by exploiting pure natural language.

Document / Passage

Caren watched TV last night. There was a guy playing tennis. Caren did not know who he is. He was wearing white shirts …

Question

What was he doing?

Answer

Playing tennis

Visual QA

Several questions require context outside of pure language.

Question

What was he doing?

Answer

Playing tennis

Visual Question Answering

Visual QA Datasets

Recently, there are a number of visual QA datasets have sprung up. Some of the more popular ones include:

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

Type #1: Real images
VQA v2.0 (Goyal et al. 2017)

Where is the child sitting?
fridge   arms

How many children are in the bed?
2   1

Type #2: Semantic reasoning
CLEVER (Johnson et al. 2016)

Q: Are there an **equal number** of **large things** and **metal spheres**?
Q: **What size** is the **cylinder that is left of** the **brown metal** thing **that is left of** the **big sphere**?
Q: There is a **sphere** with the **same size as** the **metal cube**; is it **made of the same material as** the **small red sphere**?
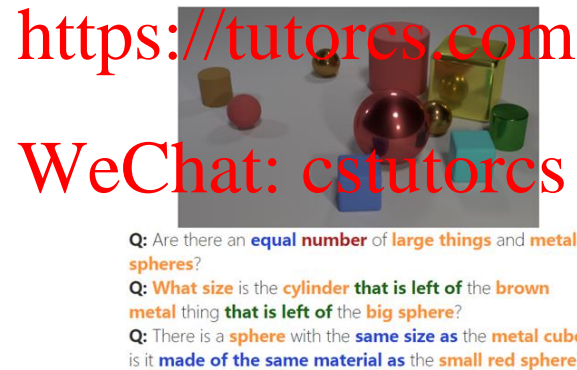
Type #3: Combined
GQA (Hudson and Manning, 2019)

Is the bowl to the right of the green apple?
What type of fruit in the image is round?
What color is the fruit on the right side, red or green?
Is there any milk in the bowl to the left of the apple?

Visual Question Answering

How does it work?

2. Context is a single picture using convolutional neural network (CNN)

4096 output units from last hidden layer
(VGGNet, Normalized)

1024

Convolution Layer
+ Non-Linearity

Convolution Layer
Pooling Layer + Non-Linearity Pooling Layer

Fully-Connected

2×2×512 LSTM

1024

1024 1000 1000

3. Single word Answer

Point-wise Fully-Connected Softmax
multiplication

"2"

Fully-Connected

"How   many   horses   are   in   this   image?"
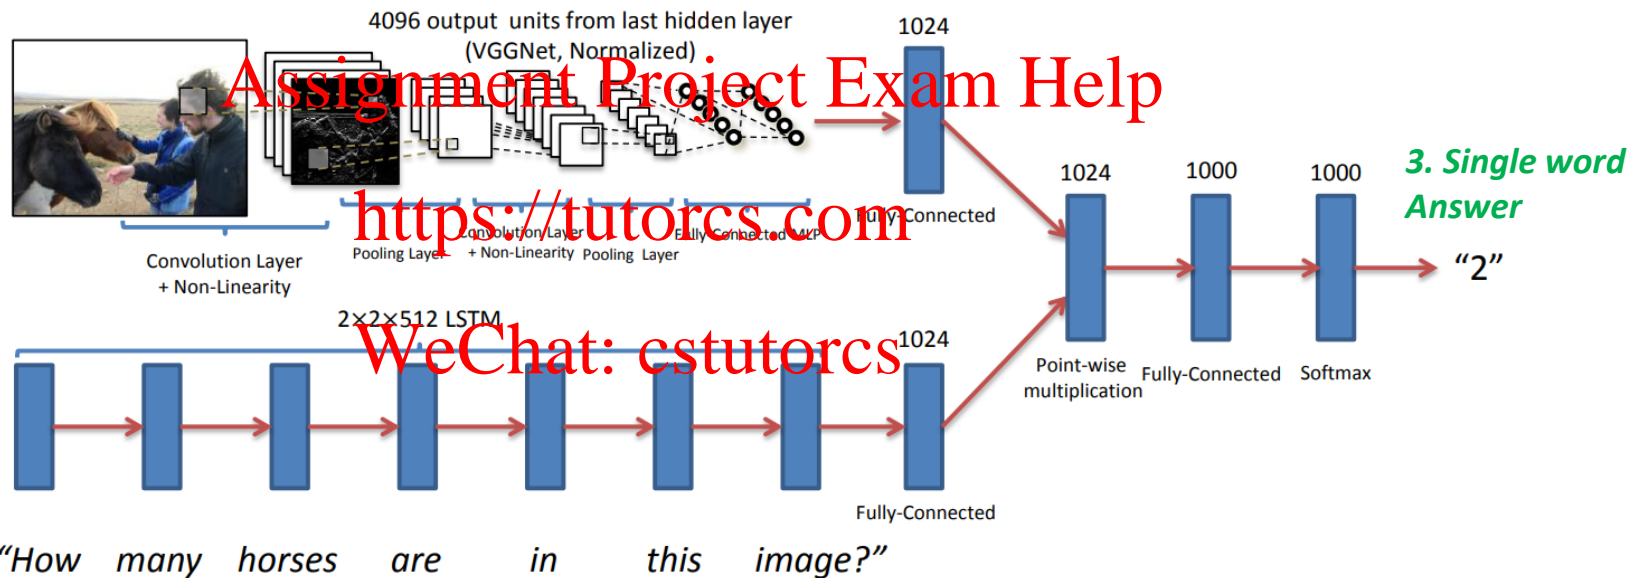
1. Encode sentence with sequence models

How does it work?

The idea of Visual QA is exactly same as reading comprehension-oriented.

*Why can't we use **Attention** then? (Yang et al. 2015)*

Answer: dogs
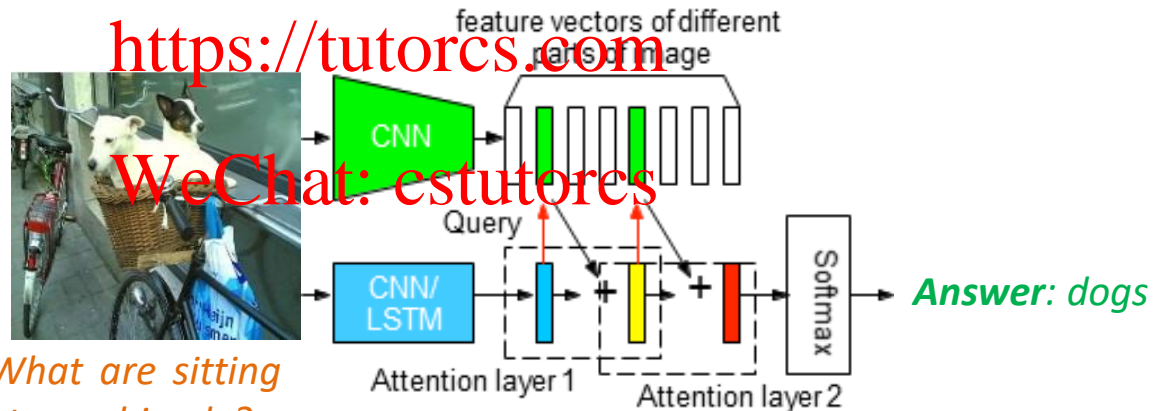
***Question**: What are sitting in the basket on a bicycle?*

Yang et al. (2015): Stacked Attention Networks for Image Question Answering

Visual Question Answering

Visual QA with Attention

Let's try some example. VisualQA (http://vqa.cloudcv.org/)

(a) What are pulling a man on a wagon down on dirt road?
Answer: horses   Prediction: horses

(b) What is the color of the box?
Answer: red Prediction: red

(c) What next to the large umbrella attached to a table?
Answer: trees Prediction: trees

(d) How many people are going up the mountain with walking sticks?
Answer: four Prediction: four

(e) What is sitting on the handle bar of a bicycle?
Answer: bird Prediction: bird

(f) What is the color of the horns?
Answer: red Prediction: red

Original Image   First Attention Layer   Second Attention Layer   Original Image   First Attention Layer   Second Attention Layer

Visual Question Answering

Visual QA with Attention (Usyd NLP Group 2020)
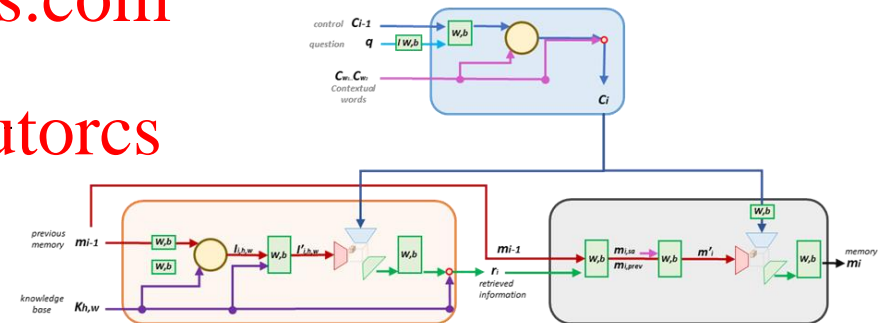
Visual QA with Attention (Usyd NLP Group 2020)
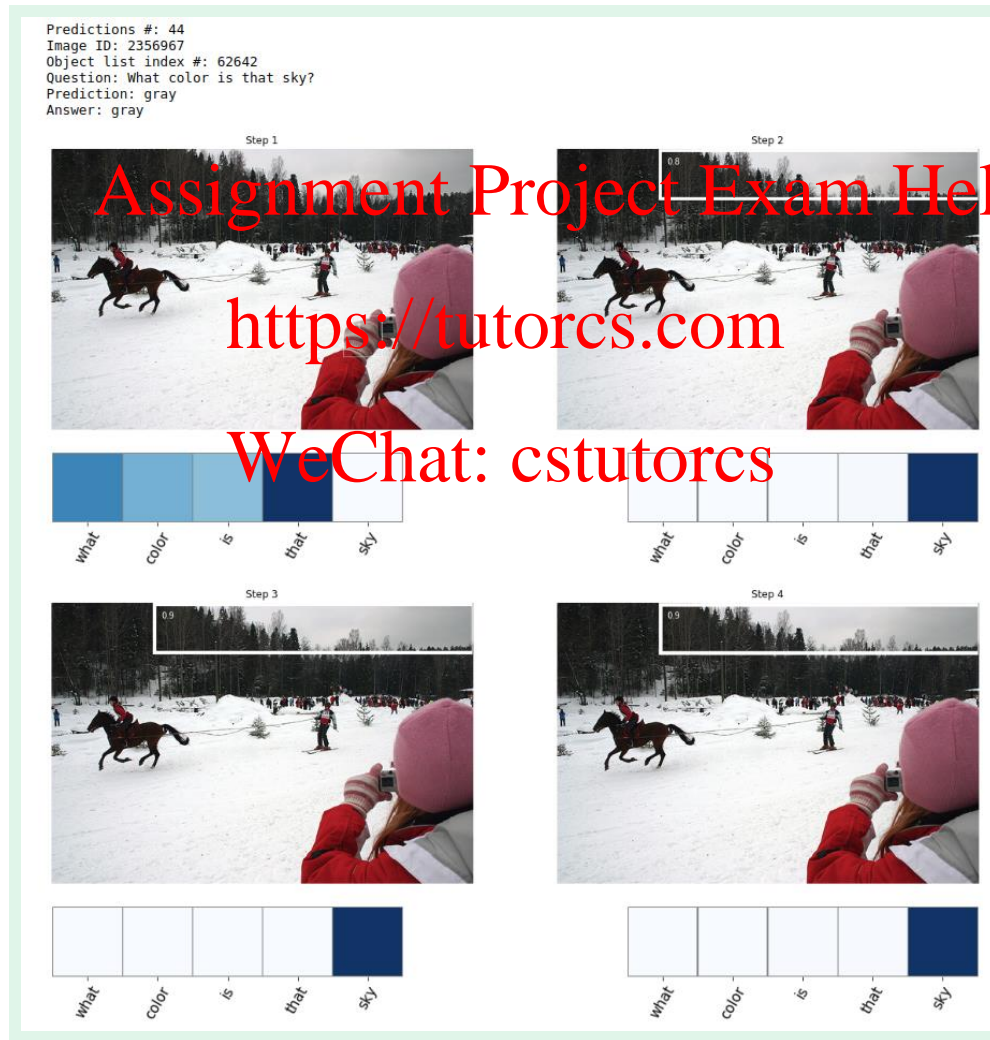
Let's have a look at the demo!!

Visual Question Answering

THE UNIVERSITY OF SYDNEY

Visual QA with Attention (Usyd NLP Group 2020)

Testing Result

Assignment Project Exam Help

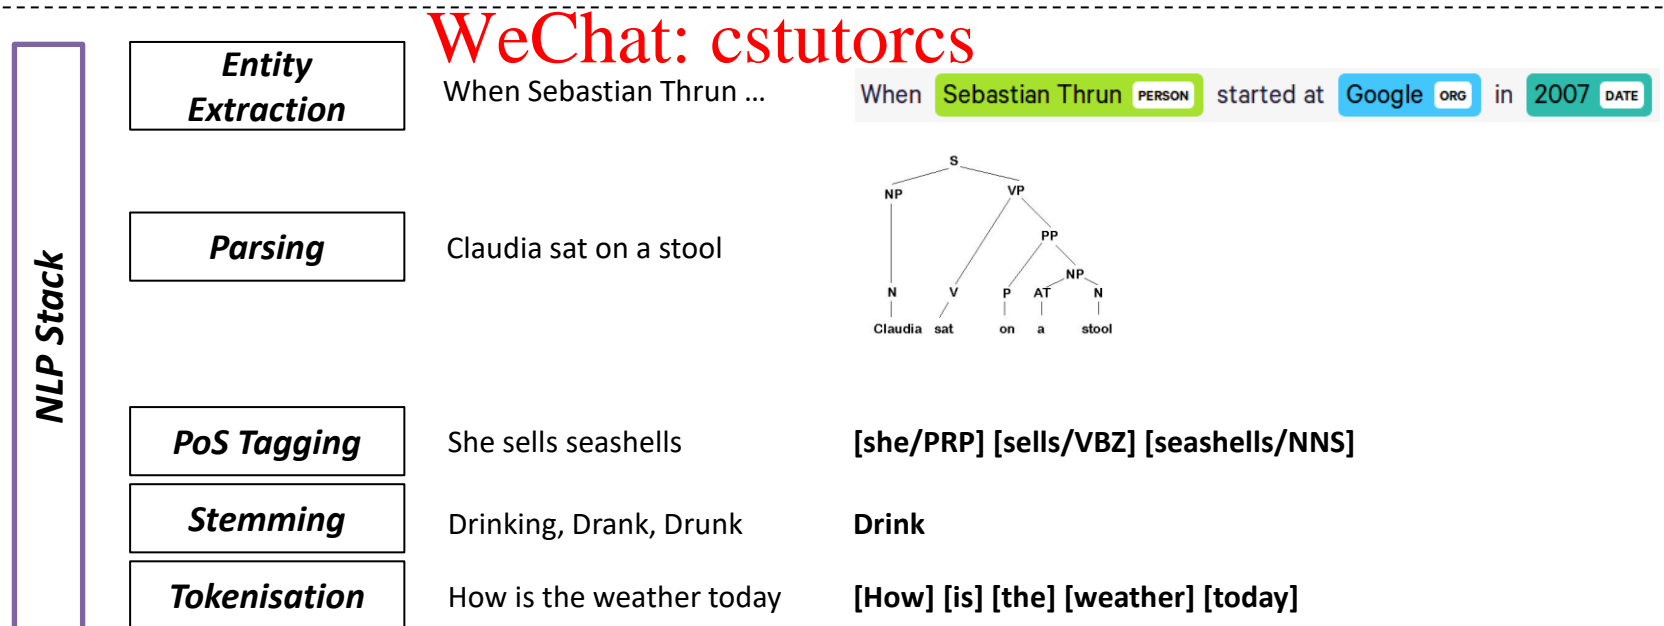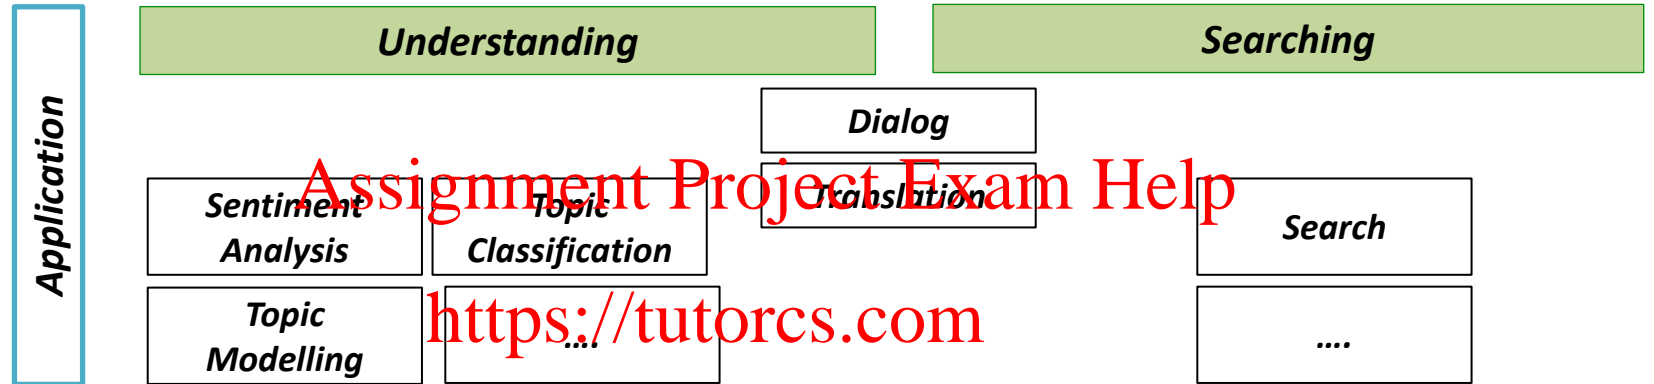*Comparison between Model Performance among State-of-the-Arts Model (GQA dataset)*

https://tutorcs.com

WeChat: cstutorcs

There is no reason to limit to just IR-based or Knowledge-based

Using multiple information sources? IBM Watson!

THE UNIVERSITY OF SYDNEY

The purpose of Natural Language Processing: Overview

Application

Understanding	Searching

Dialog

Sentiment Analysis

Topic Classification

Translation

Search

Topic Modelling

....

....

Assignment Project Exam Help

https://tutorcs.com

WeChat: cstutorcs

NLP Stack

Entity Extraction

When Sebastian Thrun ...

When Sebastian Thrun PERSON started at Google ORG in 2007 DATE

Parsing

Claudia sat on a stool

PoS Tagging

She sells seashells

[she/PRP] [sells/VBZ] [seashells/NNS]

Stemming

Drinking, Drank, Drunk

Drink

Tokenisation

How is the weather today

[How] [is] [the] [weather] [today]

Reference for this lecture

* Deng, L., & Liu, Y. (Eds.). (2018). Deep Learning in Natural Language Processing. Springer.
* Rao, D., & McMahan, B. (2019). Natural Language Processing with PyTorch: Build Intelligent Language Applications Using Deep Learning. " O'Reilly Media, Inc.".
* Manning, C. D., Manning, C. D., & Schütze, H. (1999). Foundations of statistical natural language processing. MIT press.
* Manning, C 2018, Natural Language Processing with Deep Learning, lecture notes, Stanford University

* Berant, J., Chou, A., Frostig, R., & Liang, P. (2013). Semantic parsing on freebase from question-answer pairs. In Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing (pp. 1533-1544).
* Chen, D., Bolton, J., & Manning, C. D. (2016). A thorough examination of the cnn/daily mail reading comprehension task. arXiv preprint arXiv:1606.02858.
* Chen, D., Fisch, A., Weston, J., & Bordes, A. (2017). Reading wikipedia to answer open-domain questions. arXiv preprint arXiv:1704.00051.
* Seo, M., Kembhavi, A., Farhadi, A., & Hajishirzi, H. (2016). Bidirectional attention flow for machine comprehension. arXiv preprint arXiv:1611.01603.
* Yang, Z., He, X., Gao, J., Deng, L., & Smola, A. (2016). Stacked attention networks for image question answering. In Proceedings of the IEEE conference on computer vision and pattern recognition (pp. 21-29).
* Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I. (2017). Attention is all you need. In Advances in neural information processing systems (pp. 5998-6008).
* Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2018). Bert: Pre-training of deep bidirectional transformers for language understanding. arXiv preprint arXiv:1810.04805.